

Machine Learning with Java? Deeplearning4j!

Enrique Llerena Dominguez
mimacom

Platinum
Sponsor

GP GRAND
PARADE



#DevoxxPL

@ellerenad

\$whoami

- Enrique Llerena Dominguez
- Senior Software Engineer / Tech Lead @ mimacom
- Experience developing software for the finance, retail, pharma, and automotive industry
- Professional Interests: Software Architecture, Cloud Computing, Artificial Intelligence
- Free Time: Spending time with my kids, watching football, learning german and developing nice things



Twitter @ellerenad

Github @ellerenad

ellerenad@hotmail.com

ienjoysoftware.dev

Scan the QR code!



mimacom



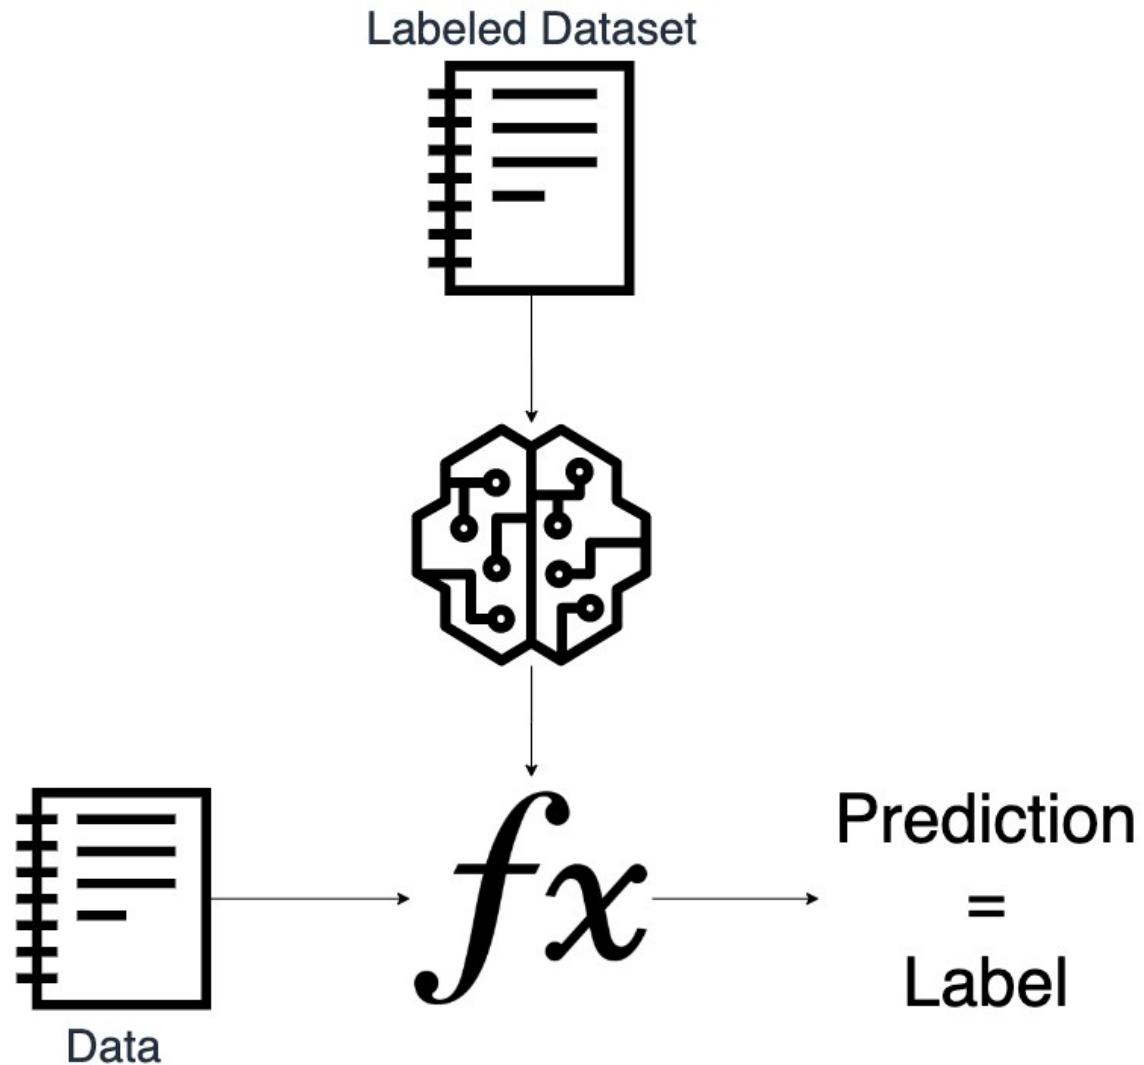
Agenda

- Machine Learning 101 + deeplearning4j
- Deeplearning4j integrations
- Hands on examples:
 - Image classification – Approach and code
 - Sentiment analysis – Approach and code

What is Machine Learning?

- Machine Learning is the study of computer algorithms that improve automatically through experience

What is Machine Learning?



What is Deeplearning4j?

Open-source deep-learning
library written for Java and
Scala



Story time!

Iris classification AKA the “Machine Learning Hello World”

Peter the biologist



Image source: Photo by [Alex](#) on [Unsplash](#)



**Gaspé
Peninsula**

Image source:

https://en.wikipedia.org/wiki/Gasp%C3%A9_Peninsula



Images source:

https://en.wikipedia.org/wiki/Iris_flower_data_set



Iris
Virginica



Iris
Versicolor



Iris
Setosa



Image source: Photo by [Blaz Erzetic](#) on [Unsplash](#)



Images source: Photo by [Ben Mullins](#) on [Unsplash](#)



Images source: Photo by [Jarritos Mexican Soda](#) on [Unsplash](#)

A photograph of a man from the waist up, wearing a dark blue plaid suit jacket over a white shirt. He is holding a bouquet of flowers consisting of small white blossoms and yellow stamens. The background is a bright, sunlit outdoor area with trees.

it would be really cool that
someone else could
classify this flower just by
providing the
measurements!

Teresa the software
engineer, Peter's
daughter



Photo by [Jarritos Mexican Soda](#) on [Unsplash](#)

hey! I can automate
that!

But first things
first



Photo by [NESA by Makers](#) on [Unsplash](#)

Resources



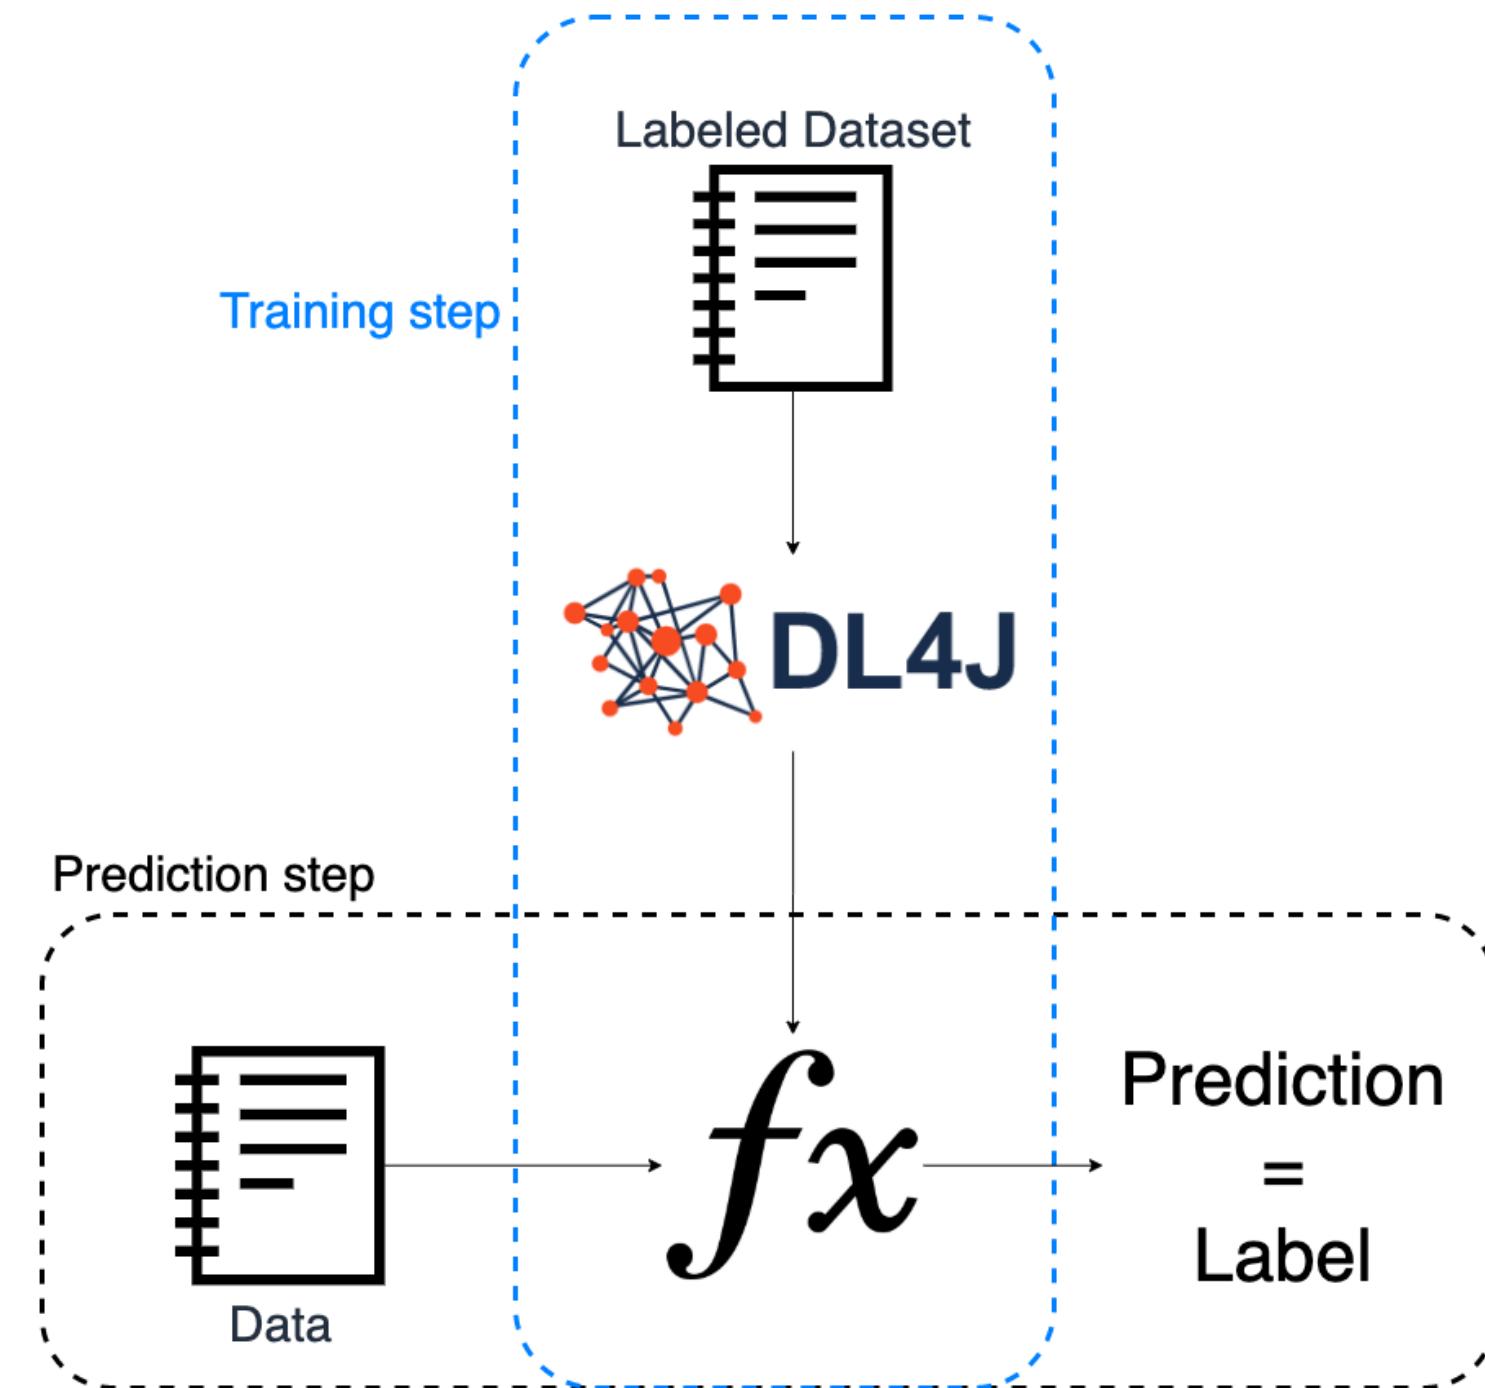
Labeled dataset



maven

Problem description

- Type: Classification
- Dataset:
 - Number of instances: 150
 - Number of attributes: 4
 - sepal length in cm
 - sepal width in cm
 - petal length in cm
 - petal width in cm
 - Number of classes: 3
 - Iris Setosa
 - Iris Versicolour
 - Iris Virginica



Training step

Load the dataset

Normalize the data

Split in train
and test
datasets

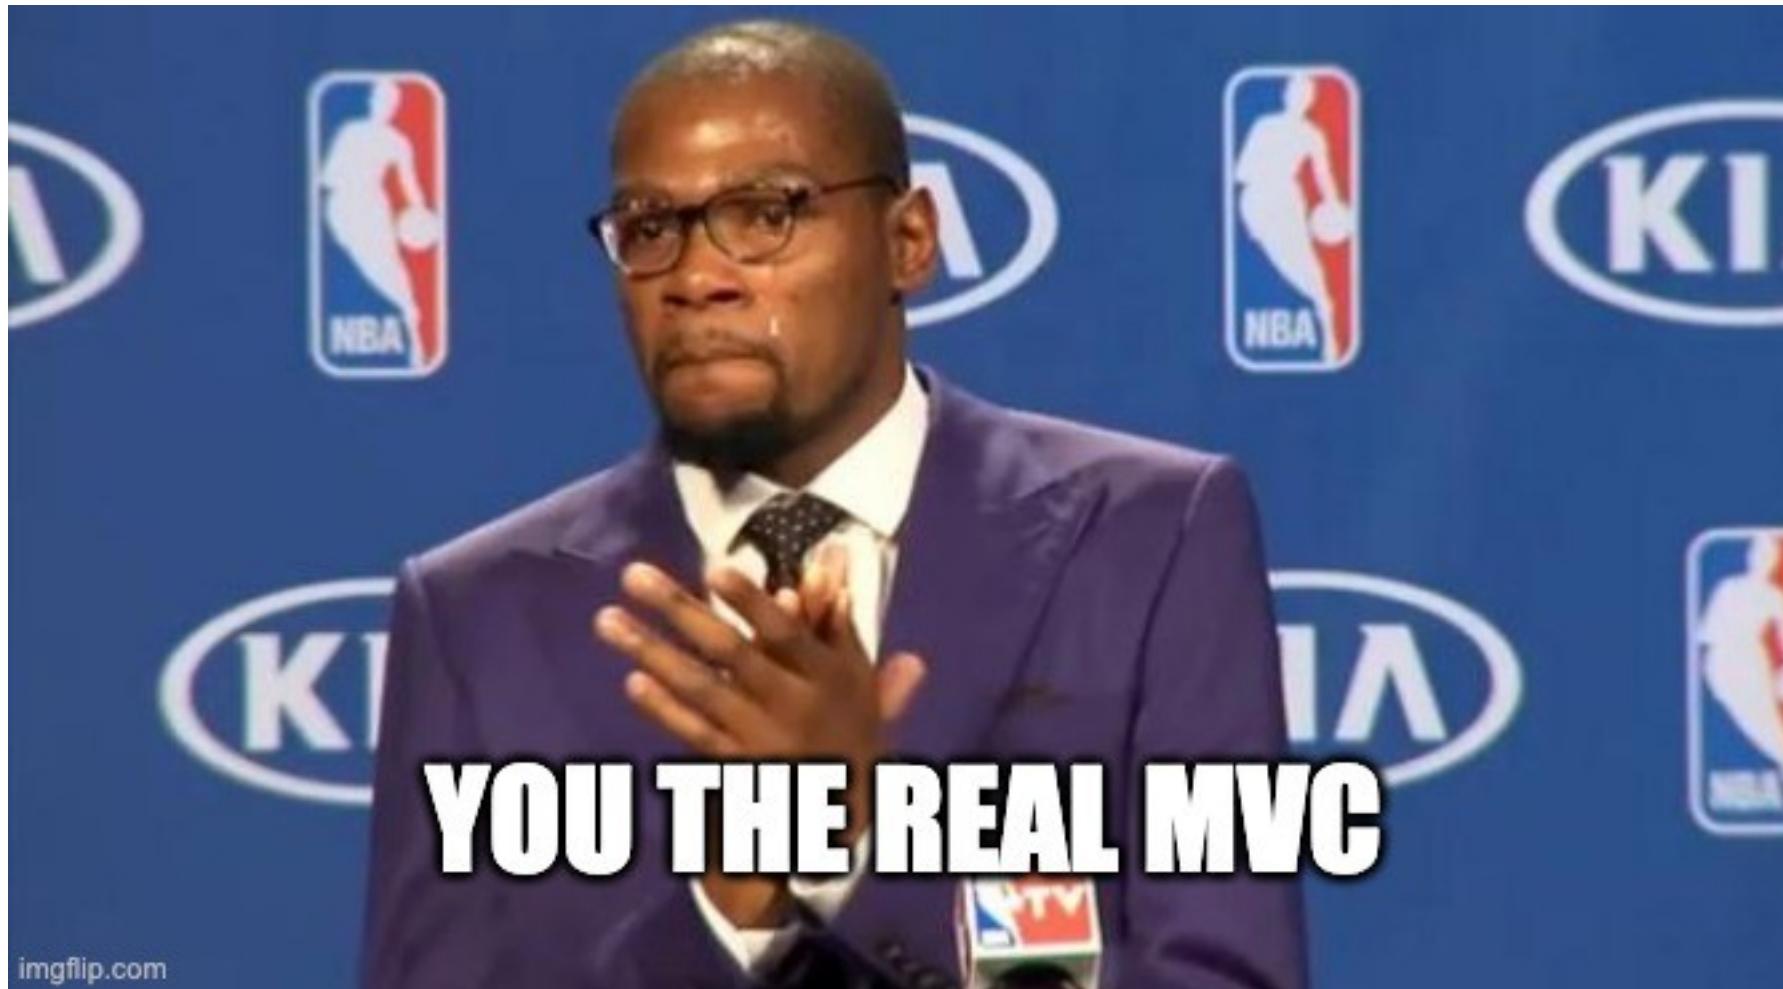
Configure
the model

Train
model

Evaluate
the model

Export the
model

Iris Classifier Trainer Most Valuable Code



Load the dataset

Normalize the data

Split in train and test datasets

Configure the model

Train model

Evaluate the model

Export the model



```
private static DataSet loadData(String path) throws IOException, InterruptedException {
    DataSet allData;
    try (RecordReader recordReader = new CSVRecordReader(0, ',')) {
        recordReader.initialize(new FileSplit(new File(path)));
        DataSetIterator iterator = new RecordReaderDataSetIterator(
                recordReader, TOTAL_LINES, LABEL_INDEX, CLASSES_COUNT);
        allData = iterator.next();
    }
    return allData;
}
```

Load the dataset

Normalize the data

Split in train and test datasets

Configure the model

Train model

Evaluate the model

Export the model



```
private static DataSet loadData(String path) throws IOException, InterruptedException {
    DataSet allData;
    try (RecordReader recordReader = new CSVRecordReader(0, ',')) {
        recordReader.initialize(new FileSplit(new File(path)));
        DataSetIterator iterator = new RecordReaderDataSetIterator(
            recordReader, TOTAL_LINES, LABEL_INDEX, CLASSES_COUNT);
        allData = iterator.next();
    }
    return allData;
}
```

Load the dataset

Normalize the data

Split in train
and test
datasets

Configure the model

Train model

Evaluate the model

Export the model



```
private static DataSet loadData(String path) throws IOException, InterruptedException {
    DataSet allData;
    try (RecordReader recordReader = new CSVRecordReader(0, ',')) {
        recordReader.initialize(new FileSplit(new File(path)));
        DataSetIterator iterator = new RecordReaderDataSetIterator(
            recordReader, TOTAL_LINES, LABEL_INDEX, CLASSES_COUNT);
        allData = iterator.next();
    }
    return allData;
}
```

```
new RecordReaderDataSetIterator(  
    recordReader, TOTAL_LINES, LABEL_INDEX, CLASSES_COUNT);
```

Load the dataset

Normalize the data

Split in train and test datasets

Configure the model

Train model

Evaluate the model

Export the model



```
// Normalize the data
DataNormalization normalizer = new NormalizerStandardize();
normalizer.fit(allData); // Get stats about the data
normalizer.transform(allData); // Transform the data by applying the normalization
```

Load the dataset

Normalize the data

Split in train and test datasets

Configure the model

Train model

Evaluate the model

Export the model



```
// Split in train and test datasets
SplitTestAndTrain testAndTrain = allData.splitTestAndTrain(TRAIN_TO_TEST_RATIO); // 65%
DataSet trainingData = testAndTrain.getTrain();
DataSet testData = testAndTrain.getTest();
```

Load the dataset

Normalize the data

Split in train and test datasets

Configure the model

Train model

Evaluate the model

Export the model

```
private static MultiLayerConfiguration getMultiLayerConfiguration() {  
    return new NeuralNetConfiguration.Builder()  
        .seed(SEED)  
        .activation(Activation.TANH)  
        .weightInit(WeightInit.XAVIER)  
        .updater(new Sgd(0.1))  
        .l2(1e-4)  
        .list()  
        // The input layer must have FEATURES_COUNT = 4 nodes  
        .layer(new DenseLayer.Builder().nIn(FEATURES_COUNT).nOut(3)  
              .build())  
        .layer(new DenseLayer.Builder().nIn(3).nOut(3)  
              .build())  
        .layer(new OutputLayer.Builder(LossFunctions.LossFunction.NEGATIVELOGLIKELIHOOD)  
              .activation(Activation.SOFTMAX)  
              .nIn(3)  
              // The output layer must have CLASSES_COUNT = 3 nodes  
              .nOut(CLASSES_COUNT).build())  
        .build();  
}
```

Load the dataset

Normalize the data

Split in train
and test
datasets

Configure
the model

Train
model

Evaluate
the model

Export the
model



```
private static MultiLayerConfiguration getMultiLayerConfiguration() {  
    return new NeuralNetConfiguration.Builder()  
        .seed(SEED)  
        .activation(Activation.TANH)  
        .weightInit(WeightInit.XAVIER)  
        .updater(new Sgd(0.1))  
        .l2(1e-4)  
        .list()  
        // The input layer must have FEATURES_COUNT = 4 nodes  
        .layer(new DenseLayer.Builder().nIn(FEATURES_COUNT).nOut(3)  
              .build())  
        .layer(new DenseLayer.Builder().nIn(3).nOut(3)  
              .build())  
        .layer(new OutputLayer.Builder(LossFunctions.LossFunction.NEGATIVELOGLIKELIHOOD)  
              .activation(Activation.SOFTMAX)  
              .nIn(3)  
              // The output layer must have CLASSES_COUNT = 3 nodes  
              .nOut(CLASSES_COUNT).build())  
        .build();  
}
```

Load the dataset

Normalize the data

Split in train
and test
datasets

Configure
the model

Train
model

Evaluate
the model

Export the
model



```
private static MultiLayerConfiguration getMultiLayerConfiguration() {  
    return new NeuralNetConfiguration.Builder()  
        .seed(SEED)  
        .activation(Activation.TANH)  
        .weightInit(WeightInit.XAVIER)  
        .updater(new Sgd(0.1))  
        .l2(1e-4)  
        .list()  
        // The input layer must have FEATURES_COUNT = 4 nodes  
        .layer(new DenseLayer.Builder().nIn(FEATURES_COUNT).nOut(3)  
              .build())  
        .layer(new DenseLayer.Builder().nIn(3).nOut(3)  
              .build())  
        .layer(new OutputLayer.Builder(LossFunctions.LossFunction.NEGATIVELOGLIKELIHOOD)  
              .activation(Activation.SOFTMAX)  
              .nIn(3)  
              // The output layer must have CLASSES_COUNT = 3 nodes  
              .nOut(CLASSES_COUNT).build())  
        .build();  
}
```

sepal length	sepal width	petal length	petal width
5,10	3,50	1,40	0,20

Types
Iris Setosa
Iris Virginica
Iris Versicolor

Load the dataset

Normalize the data

Split in train and test datasets

Configure the model

Train model

Evaluate the model

Export the model



```
// Get configuration of the Neural Network
MultiLayerConfiguration configuration = getMultiLayerConfiguration();

// Train Neural Network
MultiLayerNetwork model = new MultiLayerNetwork(configuration);
model.init();
//Print score every 100 parameter updates
model.setListeners(new ScoreIterationListener(100));

// Do TRAIN_ITERATIONS = 1000 iterations to train the model
for(int x = 0; x < TRAIN_ITERATIONS; x++) {
    model.fit(trainingData);
}
```

Load the dataset

Normalize the data

Split in train and test datasets

Configure the model

Train model

Evaluate the model

Export the model



```
private static Evaluation evaluate(MultiLayerNetwork model,  
        DataSet testData) {  
    INDArray output = model.output(testData.getFeatures());  
    Evaluation eval = new Evaluation(CLASSES_COUNT); // 4 Classes  
    eval.eval(testData.getLabels(), output);  
    return eval;  
}
```

Load the dataset

Normalize the data

Split in train
and test
datasets

Configure
the model

Train
model

Evaluate
the model

Export the
model



```
=====Evaluation Metrics=====
# of classes:      3
Accuracy:          0.9245
Precision:         0.9206
Recall:            0.9167
F1 Score:          0.9163
Precision, recall & F1: macro-averaged (equally weighted avg. of 3 classes)
```

```
=====Confusion Matrix=====
  0   1   2
-----
21   0   0 |  0 = 0
  0 15   1 |  1 = 1
  0   3 13 |  2 = 2
```

Confusion matrix format: Actual (rowClass) predicted as (columnClass) N times

Load the dataset

Normalize the data

Split in train and test datasets

Configure the model

Train model

Evaluate the model

Export the model



```
// Storing the model
File locationToSaveModel = new File(outputPath + STORED_MODEL_FILENAME);
model.save(locationToSaveModel, false);

// Storing the normalizer
File locationToSaveNormalizer = new File(outputPath + STORED_NORMALIZER_FILENAME);
NormalizerSerializer.getDefault().write(normalizer, locationToSaveNormalizer);
```

Prediction step

Load the model

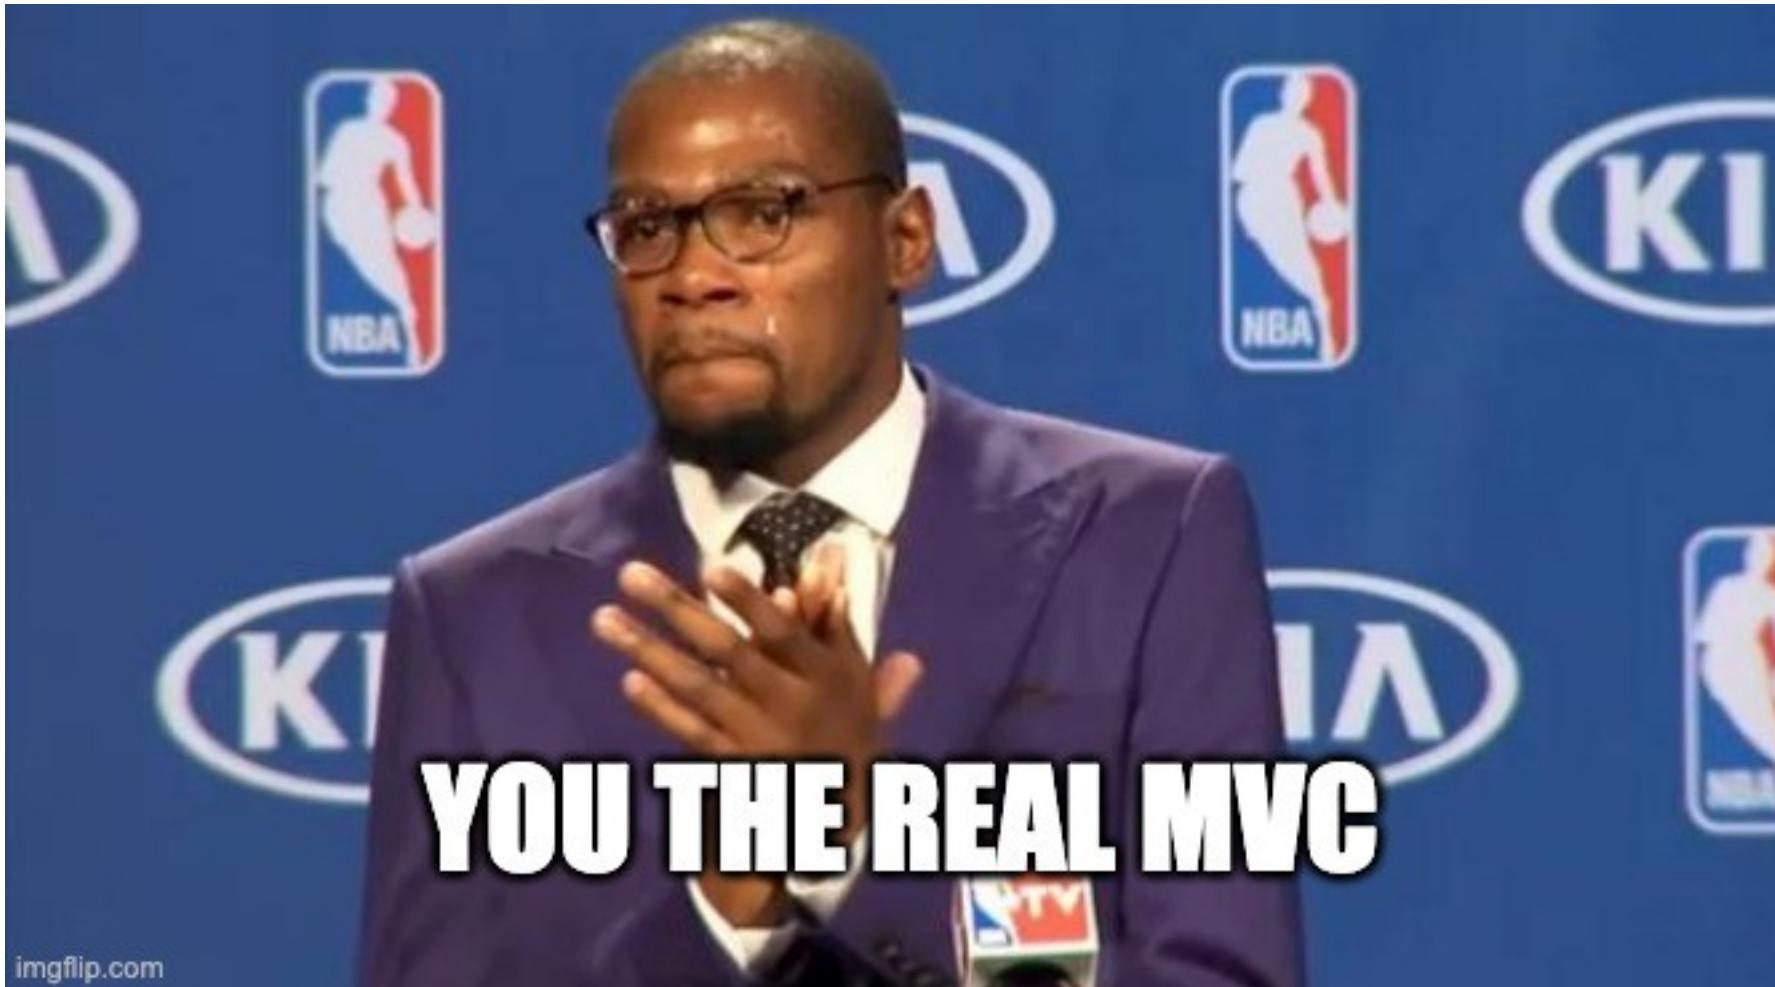
Format the data

Normalize the data

Feed the data

Get the label

Iris Classifier Predictor Most Valuable Code



Load the model

Format the data

Normalize the data

Feed the data

Get the label



```
// Load the model
File locationToSaveModel = new File(basePath + STORED_MODEL_FILENAME);
MultiLayerNetwork restoredModel = MultiLayerNetwork.load(locationToSaveModel, false);

// Load normalizer
File locationToSaveNormalizer = new File(basePath + STORED_NORMALIZER_FILENAME);
DataNormalization restoredNormalizer = normalizerSerializer.getDefault()
    .restore(locationToSaveNormalizer);
```

Load the model

Format the data

Normalize the data

Feed the data

Get the label



```
private static INDArray getArray(Iris iris) {
    float[] input = new float[FIELDS_COUNT];
    input[INDEX_SEPAL_LENGTH] = iris.getSepalLength();
    input[INDEX_SEPAL_WIDTH] = iris.getSepalWidth();
    input[INDEX_PETAL_LENGTH] = iris.getPetalLength();
    input[INDEX_PETAL_WIDTH] = iris.getPetalWidth();

    DataBuffer dataBuffer = new FloatBuffer(input);
    NDArray ndArray = new NDArray(1, FIELDS_COUNT);
    ndArray.setData(dataBuffer);
    return ndArray;
}
```

Load the model

Format the data

Normalize the data

Feed the data

Get the label



```
// Normalize the data the same way it was normalized in the training phase  
dataNormalizer.transform(indArray);  
  
// Do the prediction  
INDArray result = model.output(indArray, false);
```

Load the model

Format the data

Normalize the data

Feed the data

Get the label



```
private static int getIndexPredictedLabel(INDArray predictions) {  
    int maxIndex = 0;  
    // We should get max CLASSES_COUNT amount of predictions with probabilities.  
    for (int i = 0; i < CLASSES_COUNT; i++) {  
        if (predictions.getFloat(i) > predictions.getFloat(maxIndex)) {  
            maxIndex = i;  
        }  
    }  
    return maxIndex;  
}
```

Load the model

Format the data

Normalize the data

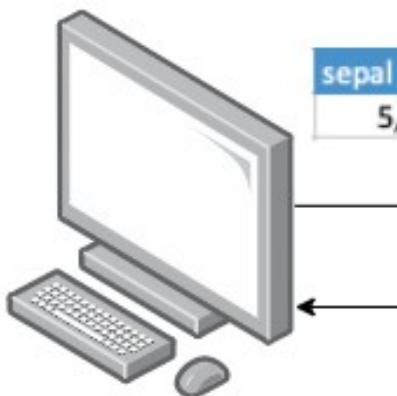
Feed the data

Get the label

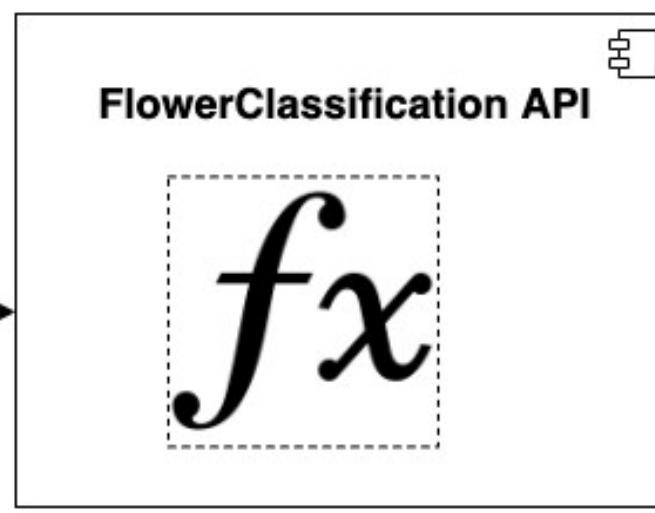


```
List<String> LABELS = Arrays.asList("Iris Setosa",
                                      "Iris Versicolour",
                                      "Iris Virginica");
int predictedLabelIndex = getIndexPredictedLabel(result);
return LABELS.get(predictedLabelIndex);
```

Prediction step



sepal length	sepal width	petal length	petal width
5,10	3,50	1,40	0,20



Iris Setosa



Now my friends can classify
these flowers just by calling
the FlowerClassification API
with the measurements! :D

Peter is
happy!

What did we learn?

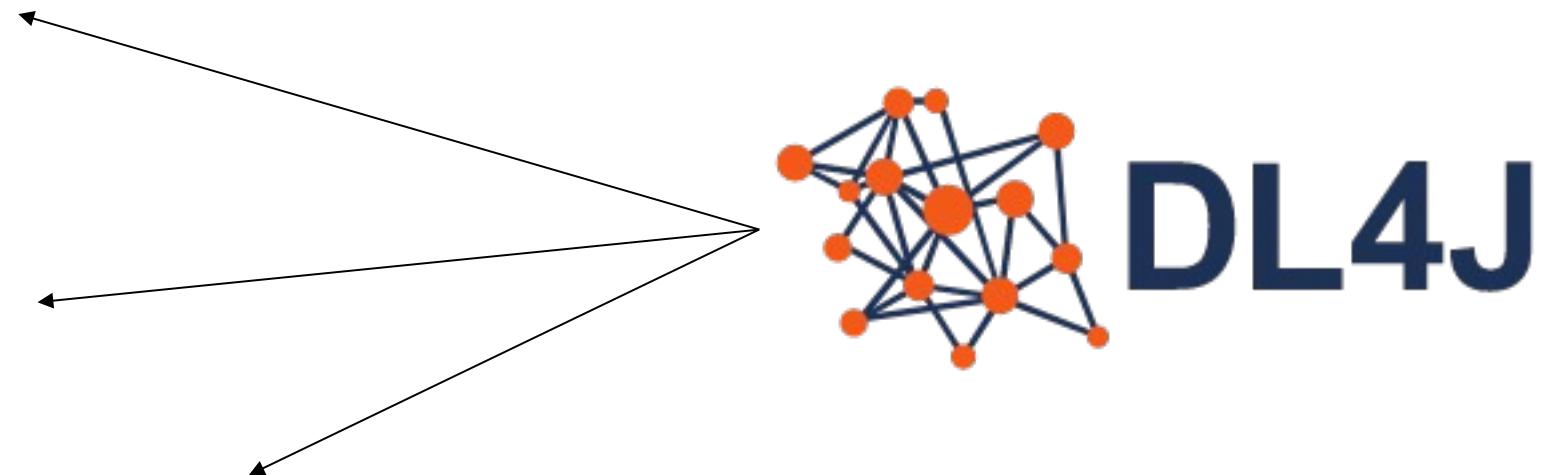
- 2 main steps
 - Training
 - Input: Labeled dataset
 - Output: Trained model
 - Prediction
 - Input: Trained model + unlabeled data
 - Output: Prediction -> in this case, of type “classification”
- How to use Deeplearning4j to train a model, export it, load it and perform a prediction

Machine Learning Approaches

- Supervised
 - Classification
 - Regression
- Unsupervised
- Reinforcement learning

Machine Learning Approaches

- Supervised
 - Classification
 - Regression
- Unsupervised
- Reinforcement learning



Machine Learning Algorithms

- Examples:
 - Support Vector Machines
 - Linear regression
 - Logistic regression
 - Naive Bayes
 - Linear discriminant analysis
 - Decision trees
 - Neural Networks (Multilayer perceptron)
 - ...



Deeplearning4j integrations



Distributed training



Deeplearning4j integrations

Parallel training



Import models

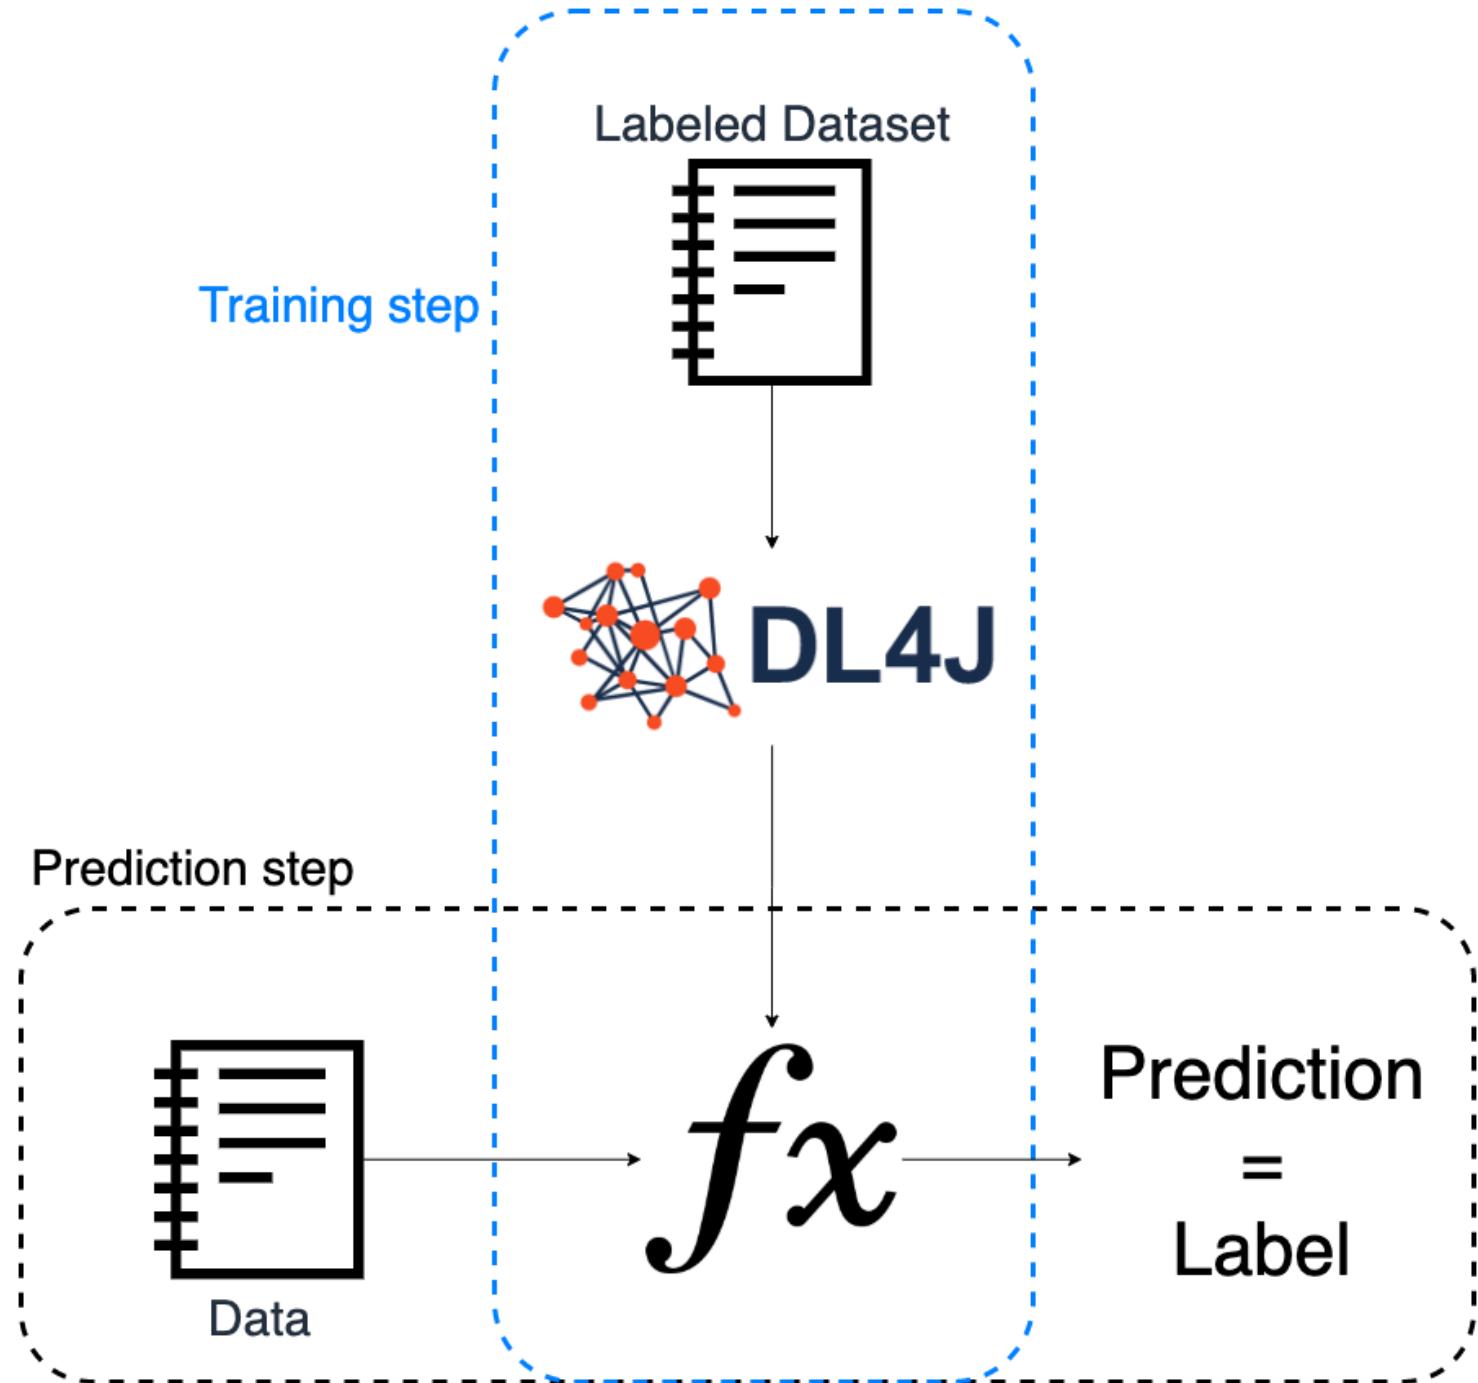


Deeplearning4j integrations



Keras is an open-source neural-network library written in Python. It is capable of running on top of TensorFlow, Microsoft Cognitive Toolkit, R, Theano, or PlaidML.

One step back



Decoupled in time!

Training step

Labeled Dataset

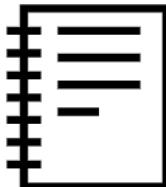


DL4J

$f\chi$

Prediction step

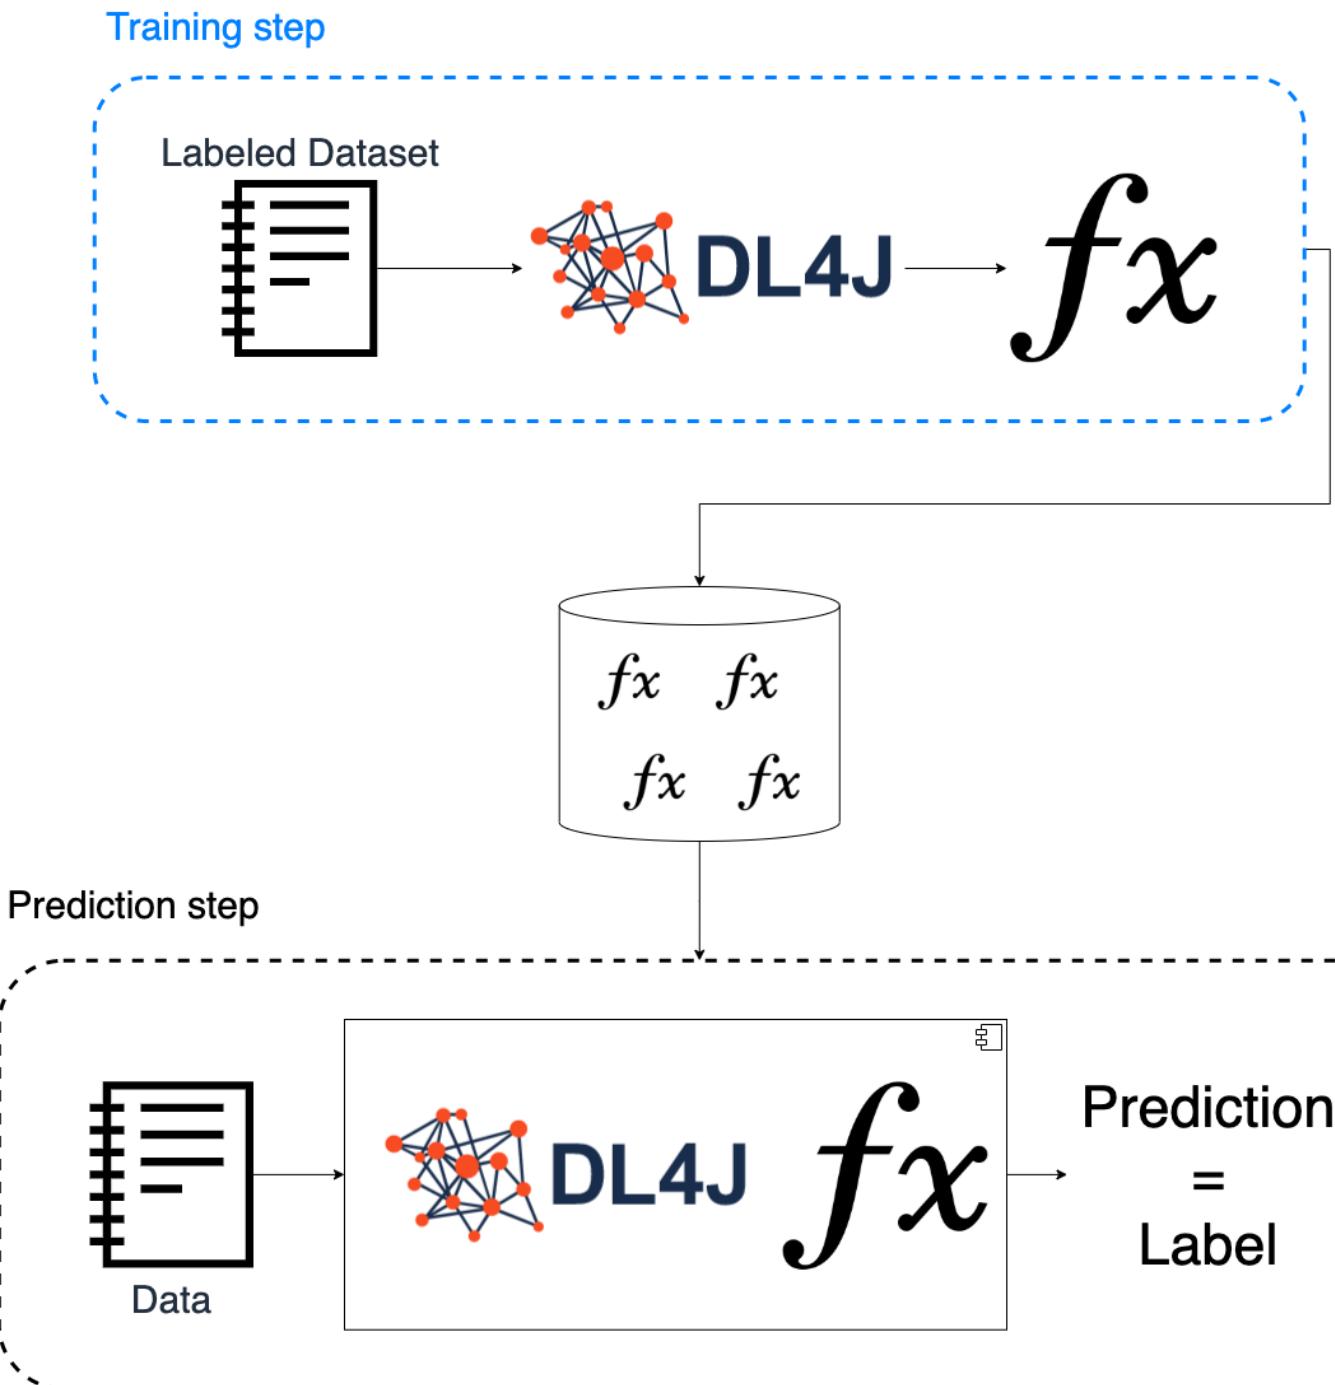
Data

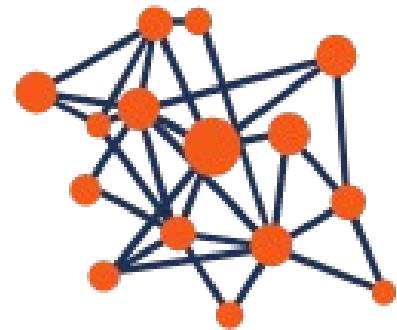


DL4J

$f\chi$

Prediction
= Label





DL4J

+



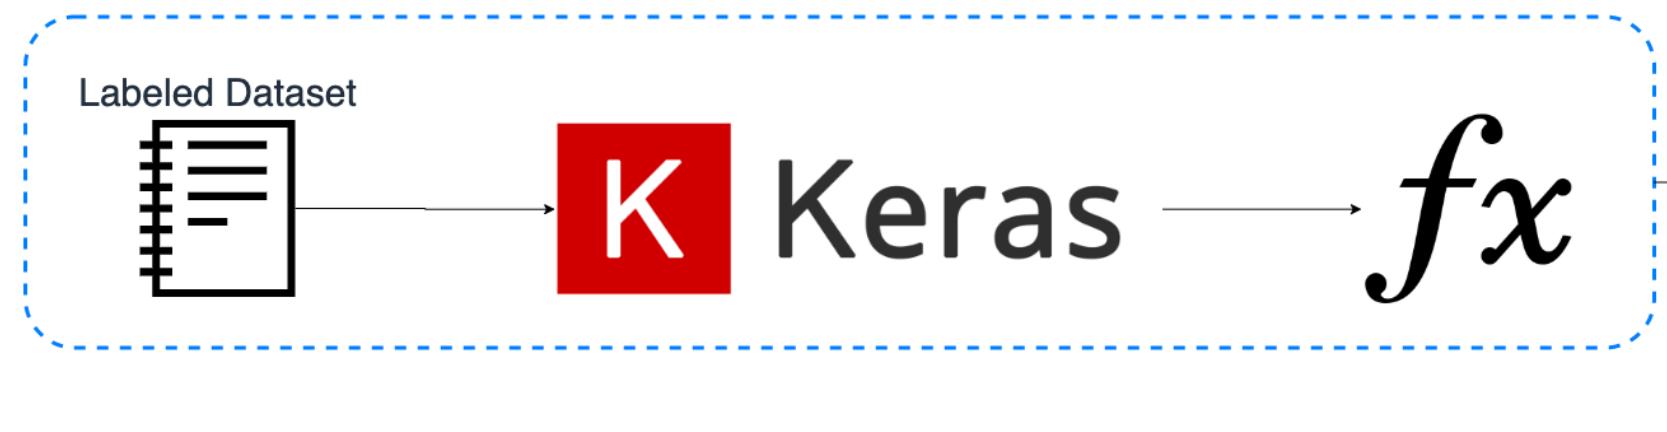
Keras

=

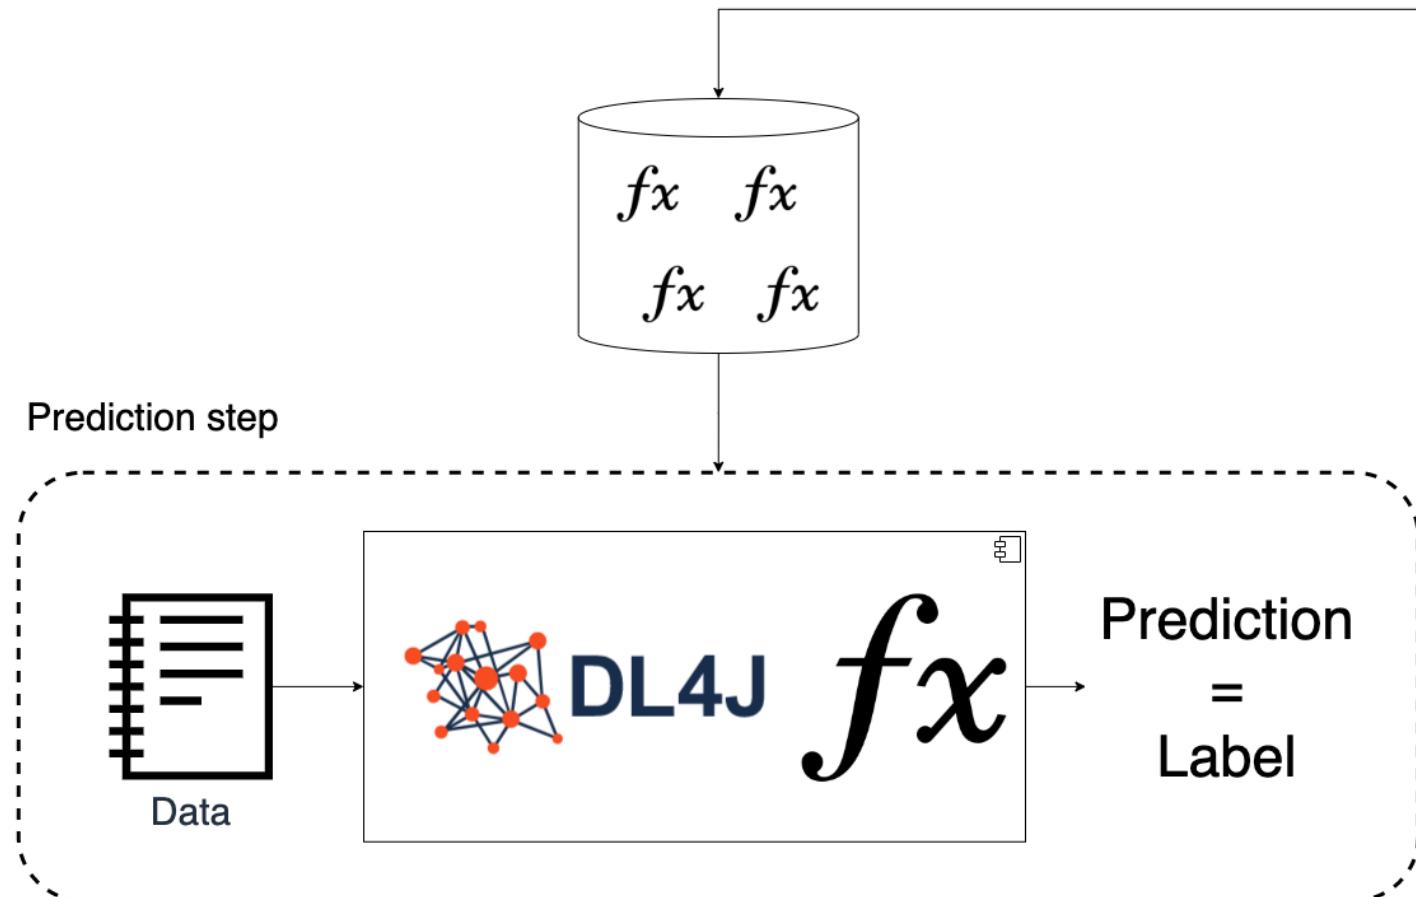
Flexibility



Training step



Prediction step



Deeplearning4j integrations

- Enables the integration of people with different technical profiles
- Opens the door to more pretrained models

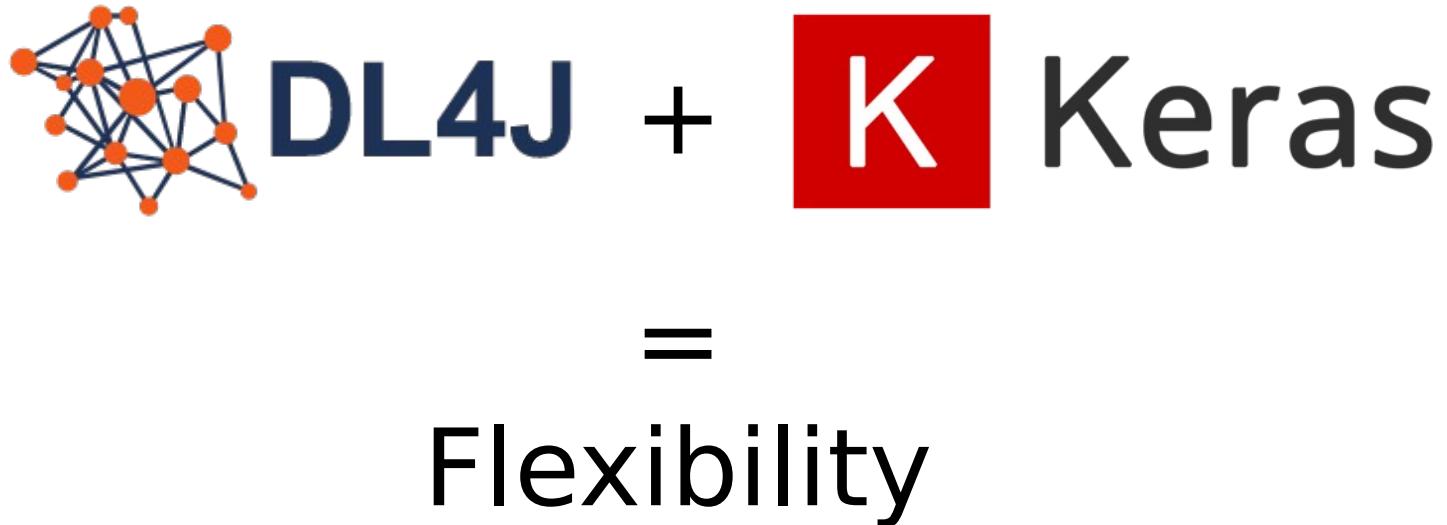


Image Classification

- Approach:
 - Deeplearning4j's Model Zoo
 - OpenCV



AlexDBlack Refactor packages to fix split package issues (#411) ...

..



helper

Refactor packages to fix split package issues (#411)



AlexNet.java

Fixing issues from Sonar report (#391)



Darknet19.java

Eclipse Migration Initial Commit



FaceNetNN4Small2.java

Eclipse Migration Initial Commit



InceptionResNetV1.java

Eclipse Migration Initial Commit



LeNet.java

Eclipse Migration Initial Commit



NASNet.java

Refactor packages to fix split package issues (#411)



ResNet50.java

Eclipse Migration Initial Commit



SimpleCNN.java

Eclipse Migration Initial Commit



SqueezeNet.java

Eclipse Migration Initial Commit



TextGenerationLSTM.java

Eclipse Migration Initial Commit



TinyYOLO.java

Eclipse Migration Initial Commit



UNet.java

Various SameDiff fixes (#21)



VGG16.java

Eclipse Migration Initial Commit



VGG19.java

Eclipse Migration Initial Commit



Xception.java

Eclipse Migration Initial Commit



YOLO2.java

Eclipse Migration Initial Commit

Deeplearning4j's Model Zoo



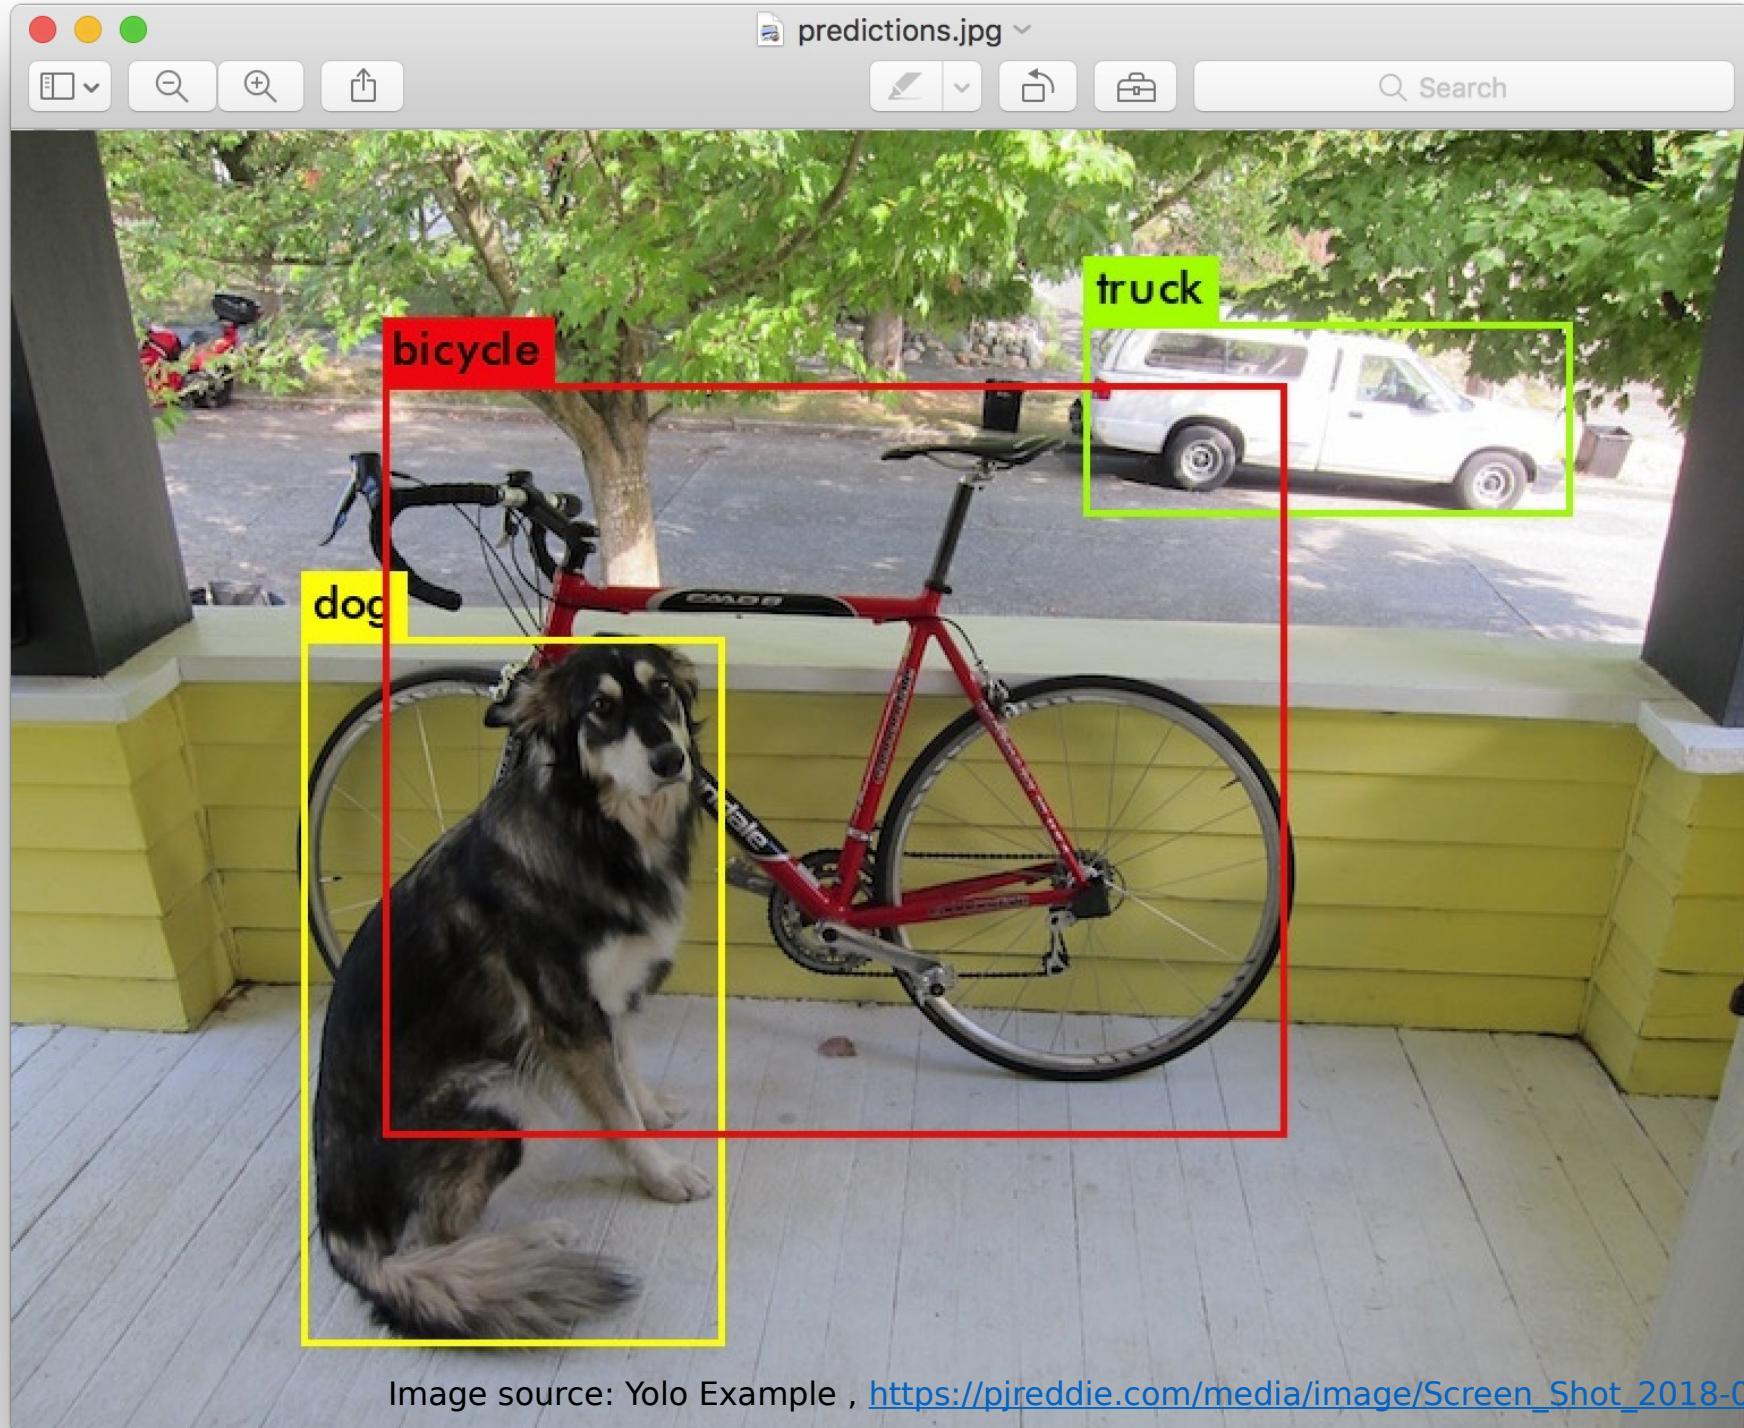


Image source: Yolo Example , https://pjreddie.com/media/image/Screen_Shot_2018-03-24_at_10.48.42_PM.png

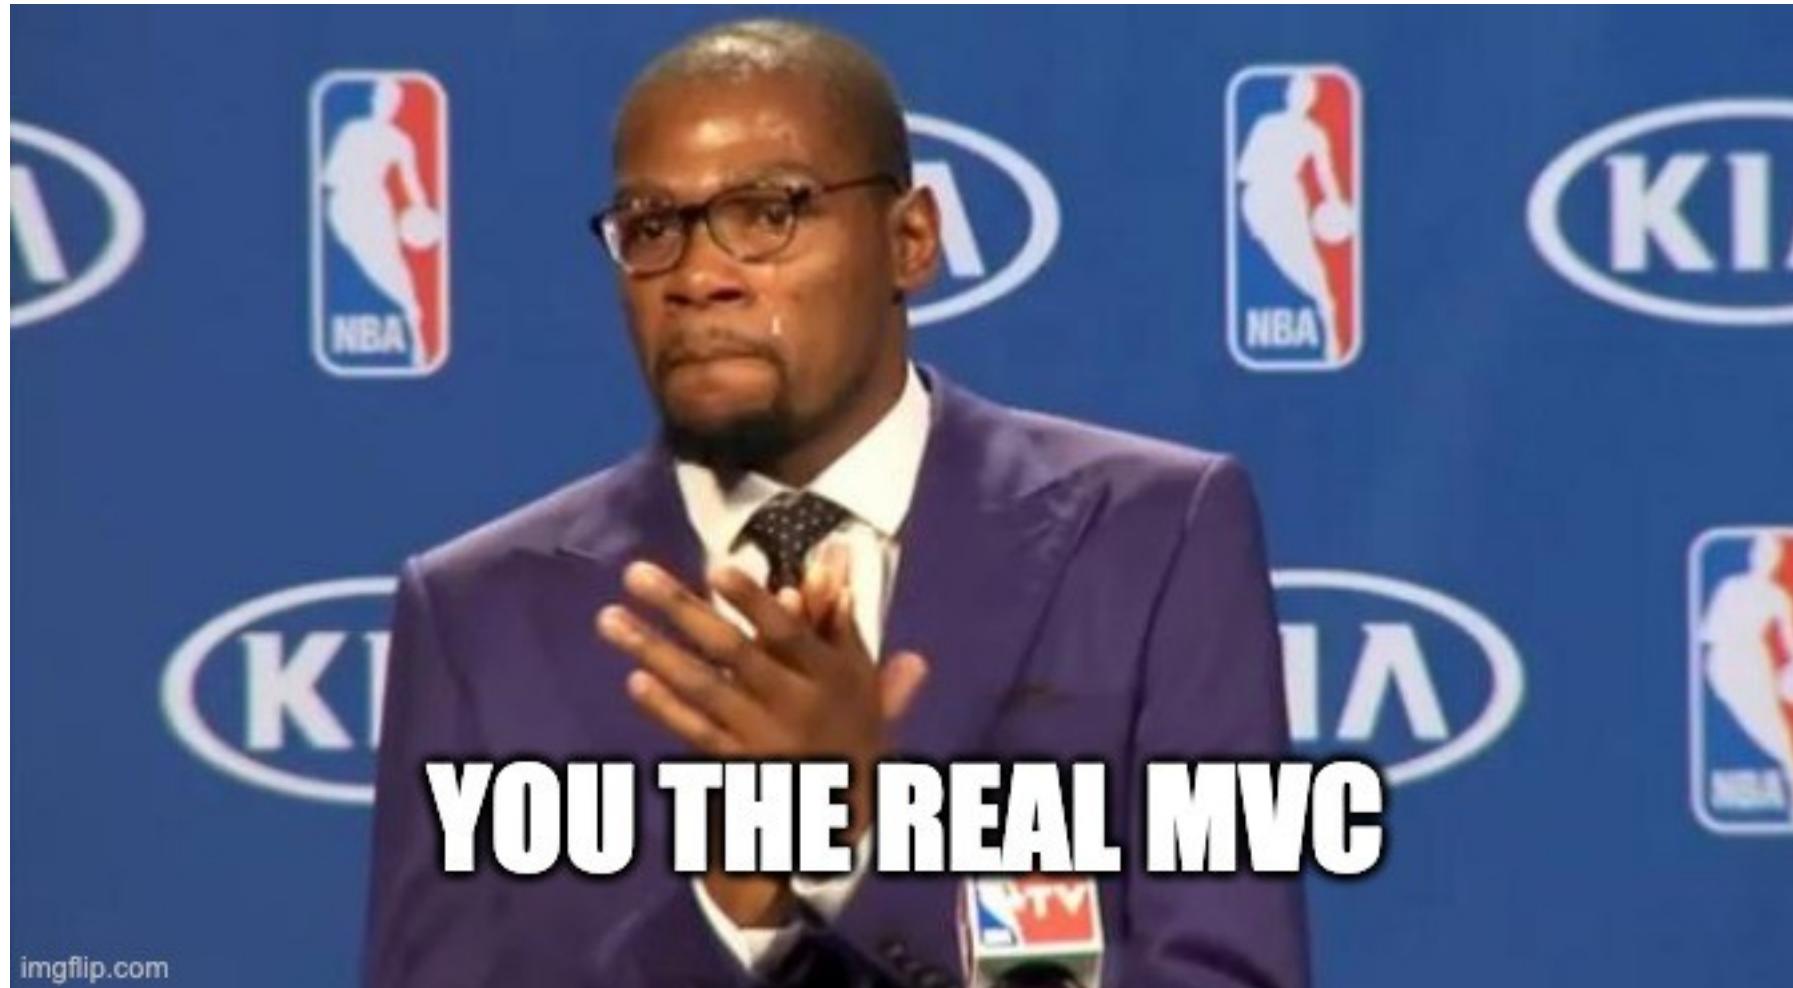
YOLOv2 Available Labels

person	fire hydrant	elephant	skis	wine glass	broccoli	diningtable	toaster
bicycle	stop sign	bear	snowboard	cup	carrot	toilet	sink
car	parking meter	zebra	sports ball	fork	hot dog	tvmonitor	refrigerator
motorbike	bench	giraffe	kite	knife	pizza	laptop	book
aeroplane	bird	backpack	baseball bat	spoon	donut	mouse	clock
bus	cat	umbrella	baseball glove	bowl	cake	remote	vase
train	dog	handbag	skateboard	banana	chair	keyboard	scissors
truck	horse	tie	surfboard	apple	sofa	cell phone	teddy bear
boat	sheep	suitcase	tennis racket	sandwich	pottedplant	microwave	hair drier
traffic light	cow	frisbee	bottle	orange	bed	oven	toothbrush

Steps overview



Yolo2ImageClassifier - Most Valuable Code



```
public List<DetectedObject> classify(String inputImagePath, String outputImagePath) throws IOException {
    // Load the model from the zoo
    yolo2Model = YOLO2.builder().build();
    pretrainedComputationGraph = (ComputationGraph) yolo2Model.initPretrained();

    // Load the image from disk
    File fileOriginalImage = new File(inputImagePath);
    INDArray iNDArrayOriginalImage = imageLoader.asMatrix(fileOriginalImage);

    // Resize the image and change the format to match the required by YOLO2
    yoloImageLoader = new NativeImageLoader(YOLO2_WIDTH, YOLO2_HEIGHT, CHANNELS,
                                           new ColorConversionTransform(COLOR_BGR2RGB));
    Mat matResizedImage = yoloImageLoader.asMat(iNDArrayOriginalImage);

    // Scale the images, as in "normalize the pixels to be on the range from 0 to 1"
    ImagePreProcessingScaler scaler = new ImagePreProcessingScaler(0, 1);
    INDArray iNDArrayTransformedImage = yoloImageLoader.asMatrix(matResizedImage);
    scaler.transform(iNDArrayTransformedImage);

    // Perform the classification
    INDArray outputs = pretrainedComputationGraph.outputSingle(iNDArrayTransformedImage);
    List<DetectedObject> detectedObjects = YoloUtils.getPredictedObjects(
        Nd4j.create((YOLO2) yolo2Model).getPriorBoxes(),
        outputs,
        DETECTION_THRESHOLD,
        NMS_THRESHOLD);

    // Annotate the original image
    Image originalImage = imageLoader.asImageMatrix(fileOriginalImage);
    int originalWidth = originalImage.getOrigW();
    int originalHeight = originalImage.getOrigH();
    annotate(originalWidth, originalHeight, matResizedImage, detectedObjects, outputImagePath);

    return detectedObjects;
}
```

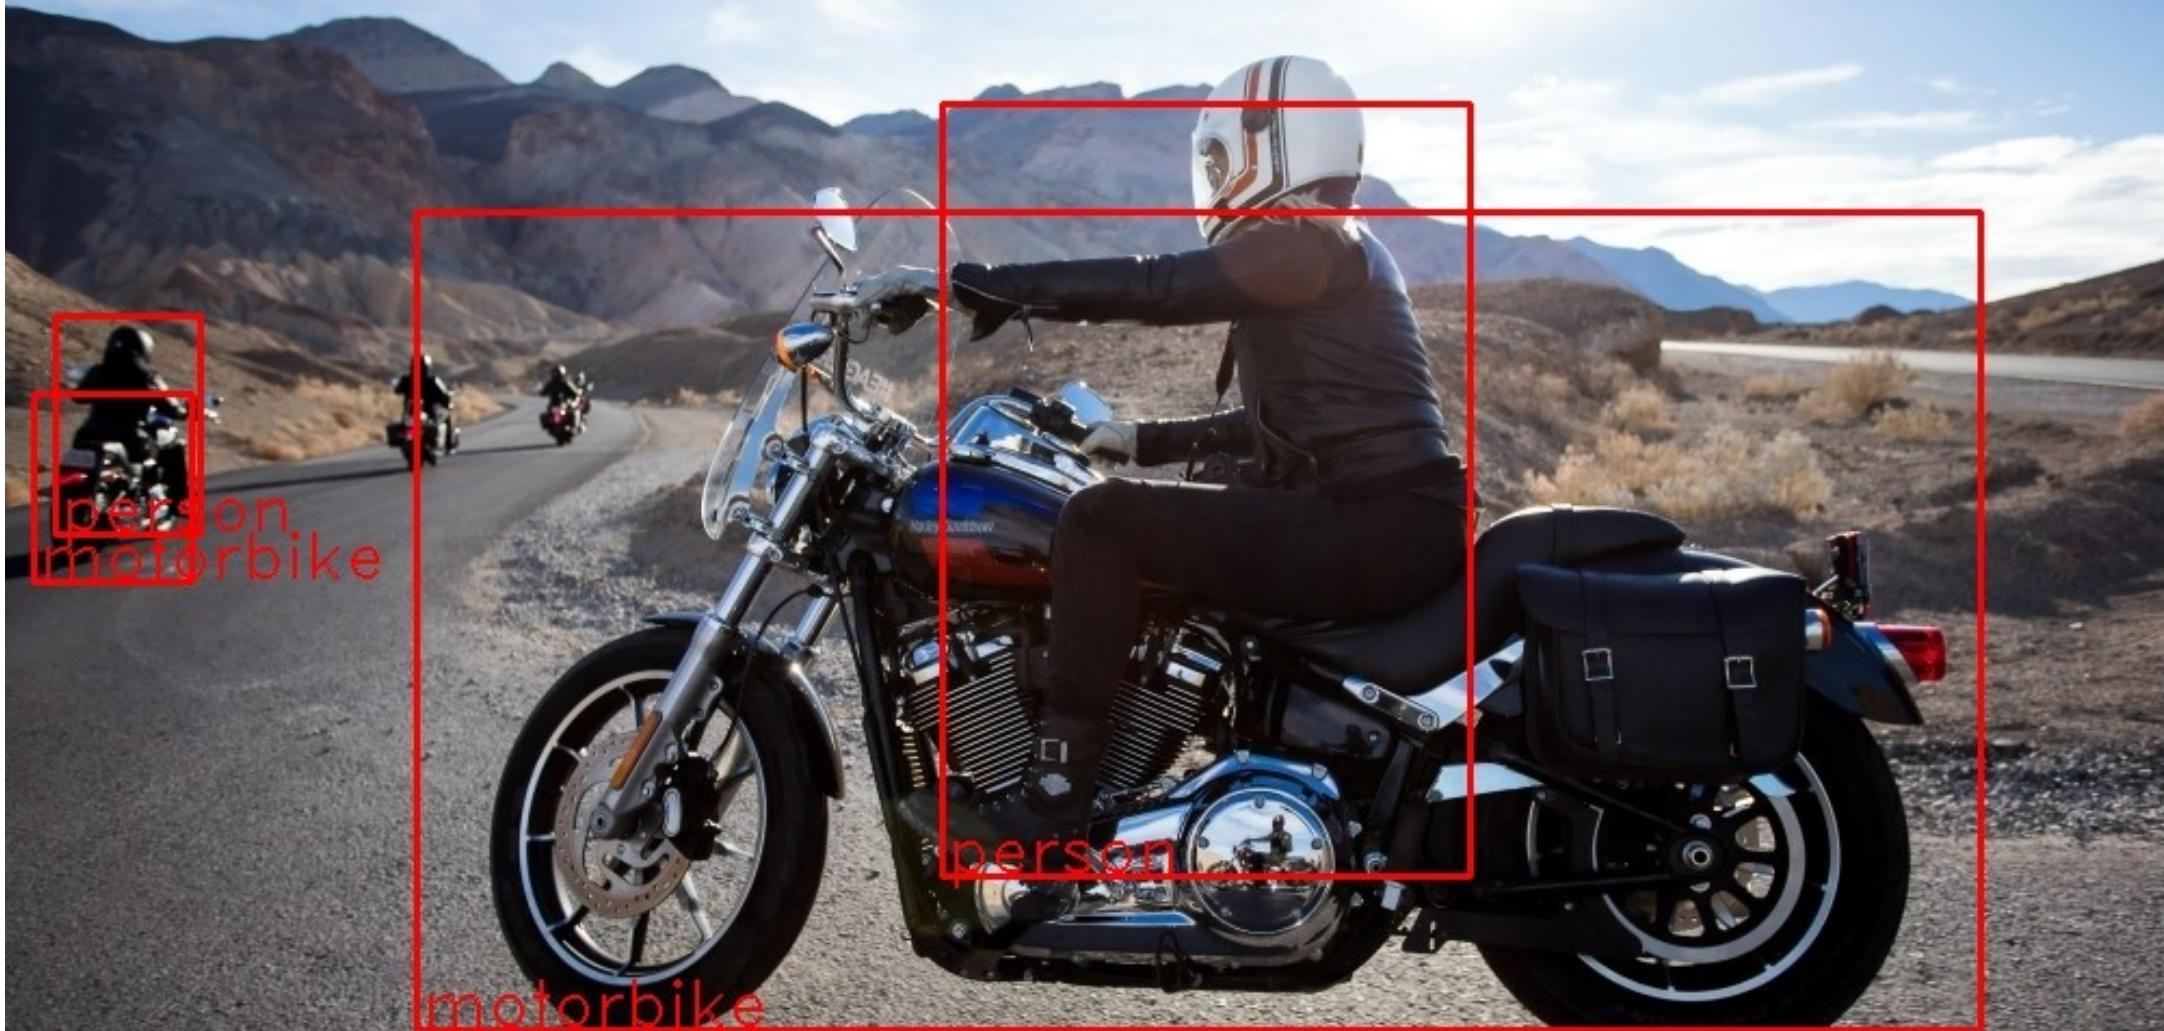


Image source: unsplash. Annotated by myself using yolo2



Image source: unsplash. Annotated by myself using yolo2



bird



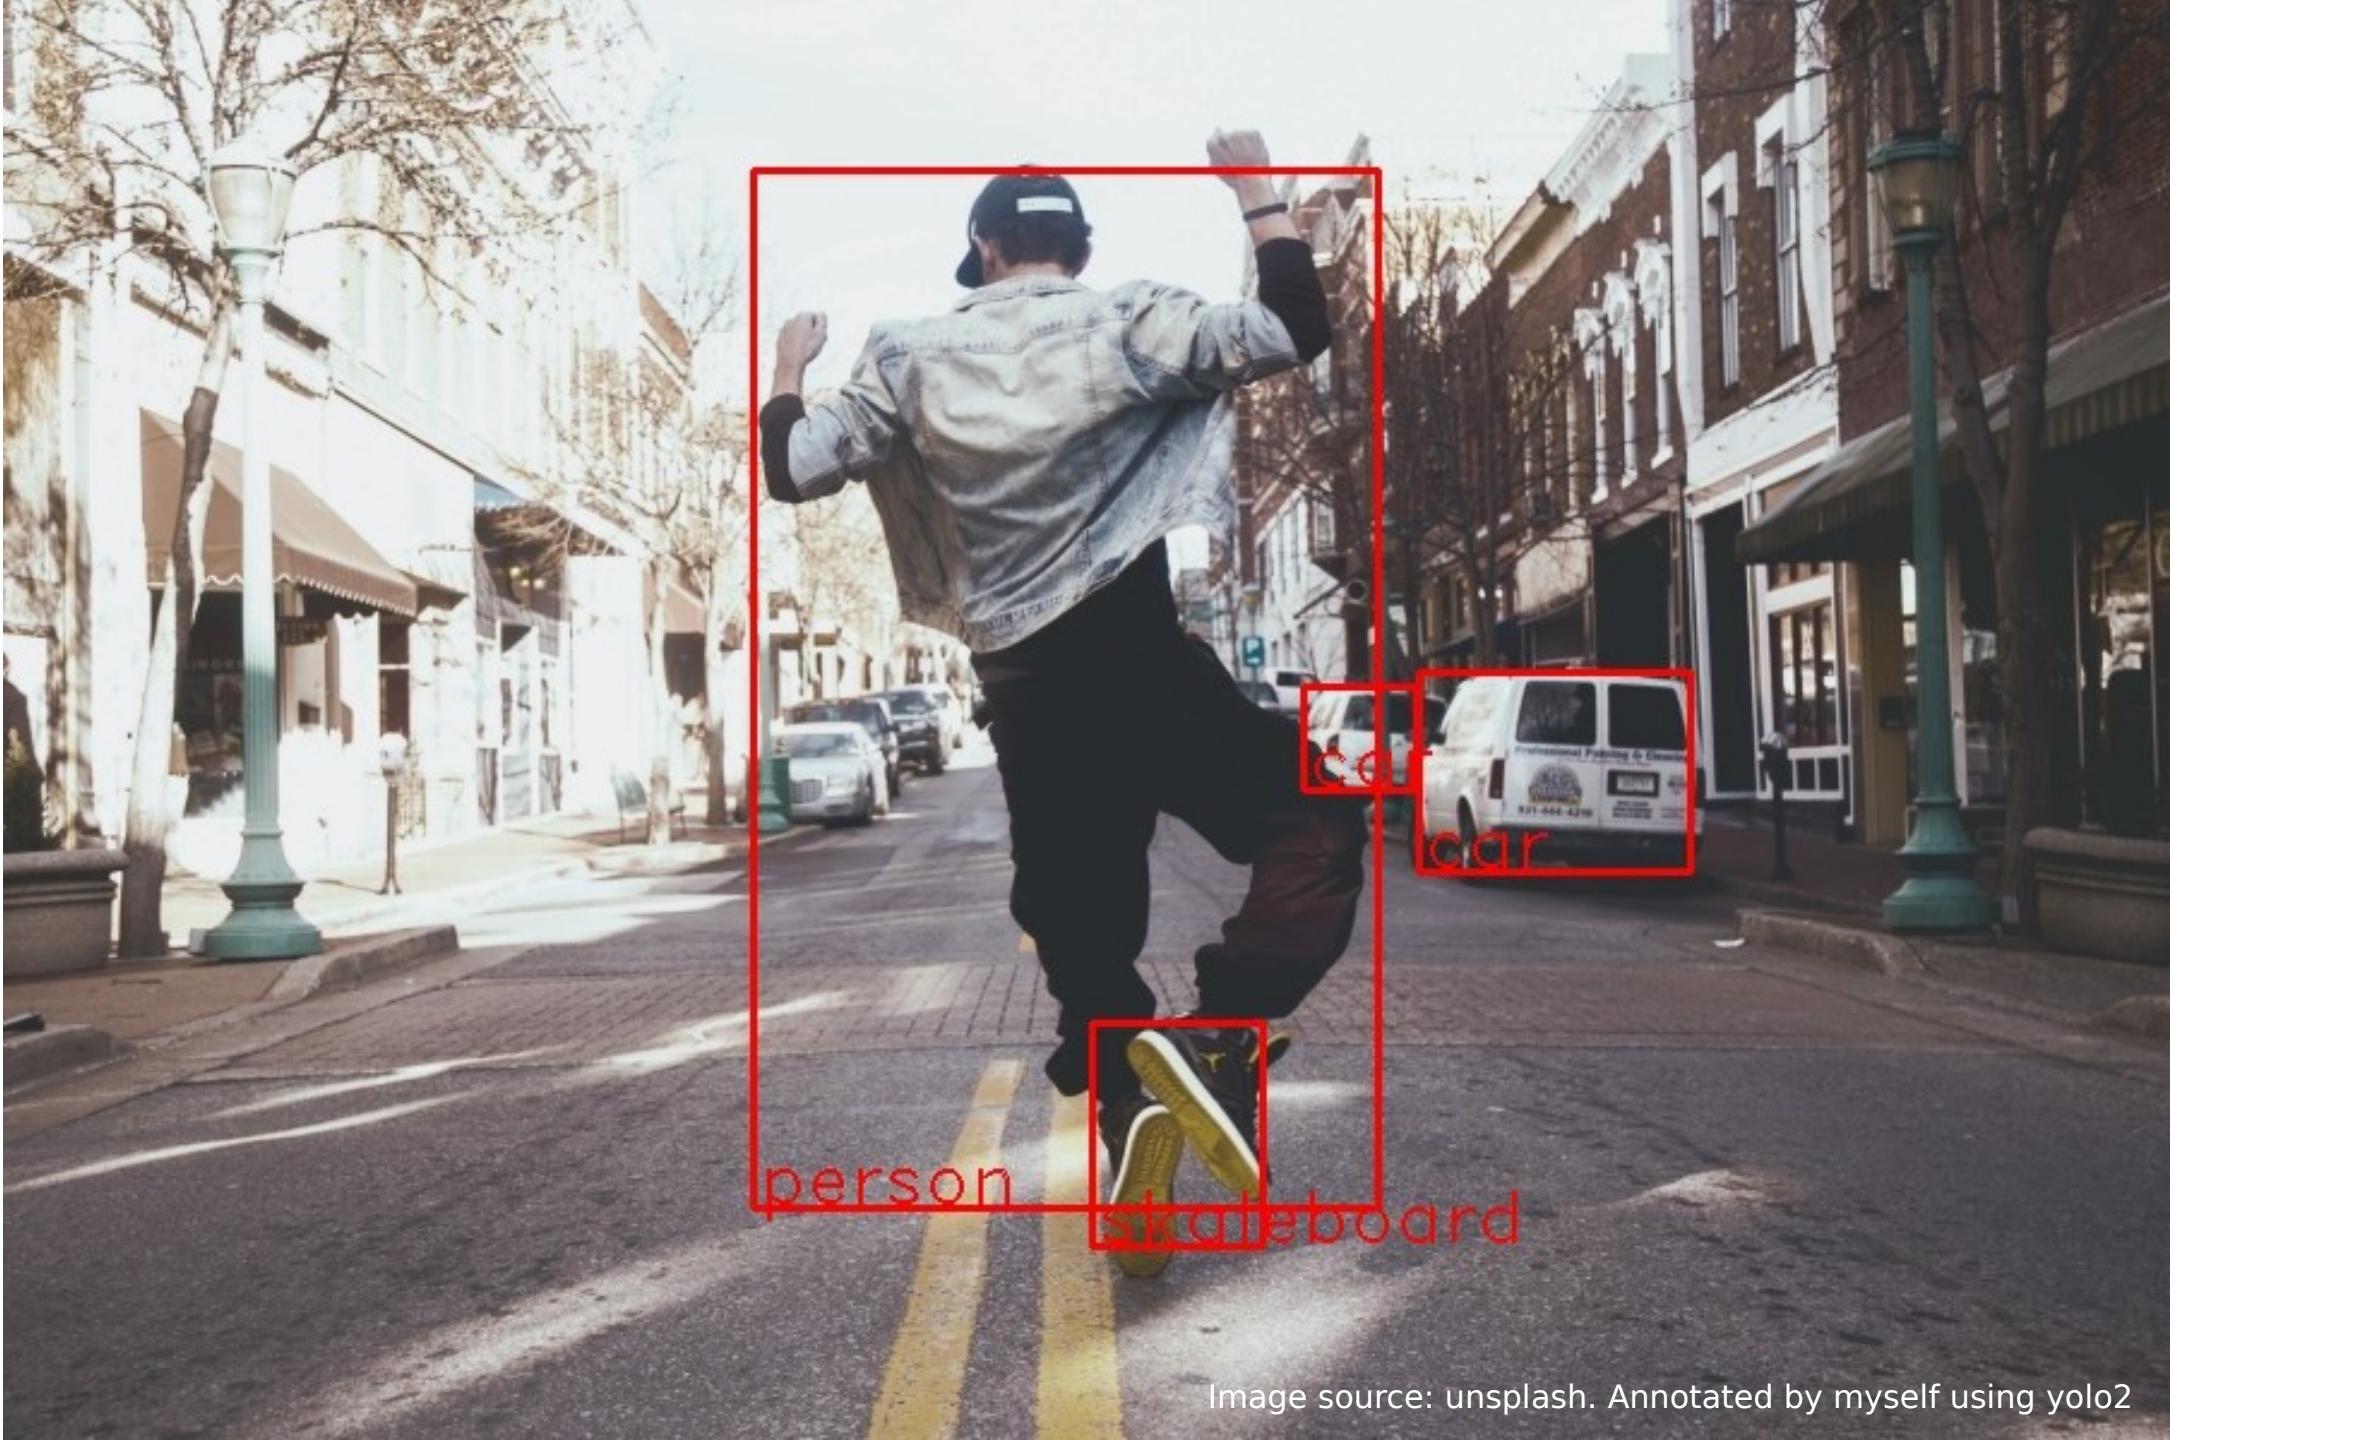


Image source: unsplash. Annotated by myself using yolo2



Image source: Deeplearning4j Animal Classification Dataset. Annotated by myself using yolo2



Image source: Deeplearning4j Animal Classification Dataset. Annotated by myself using yolo2

Sentiment Analysis

- Approach:
 - Take an example and modify it
 - deeplearning4j.examples - ImdbReviewClassificationCNN.java
 - Get a new dataset from
<https://www.kaggle.com/kazanova/sentiment140>

Sentiment140 dataset with 1.6 million tweets

kaggle.com/kazanova/sentiment140

Search

Dataset

Sentiment140 dataset with 1.6 million tweets

Sentiment analysis with tweets

Μάριος Μιχαηλίδης Kazanova • updated 3 years ago (Version 2)

Data Tasks Notebooks (121) Discussion (9) Activity Metadata Download (228 MB) New Notebook

Usability 8.8 License Other (specified in description) Tags internet, online communities, social networks, linguistics, languages

Description

Context

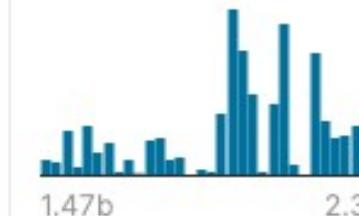
This is the sentiment140 dataset. It contains 1,600,000 tweets extracted using the twitter api . The tweets have been annotated (0 = negative, 4 = positive) and they can be used to detect sentiment .

Content

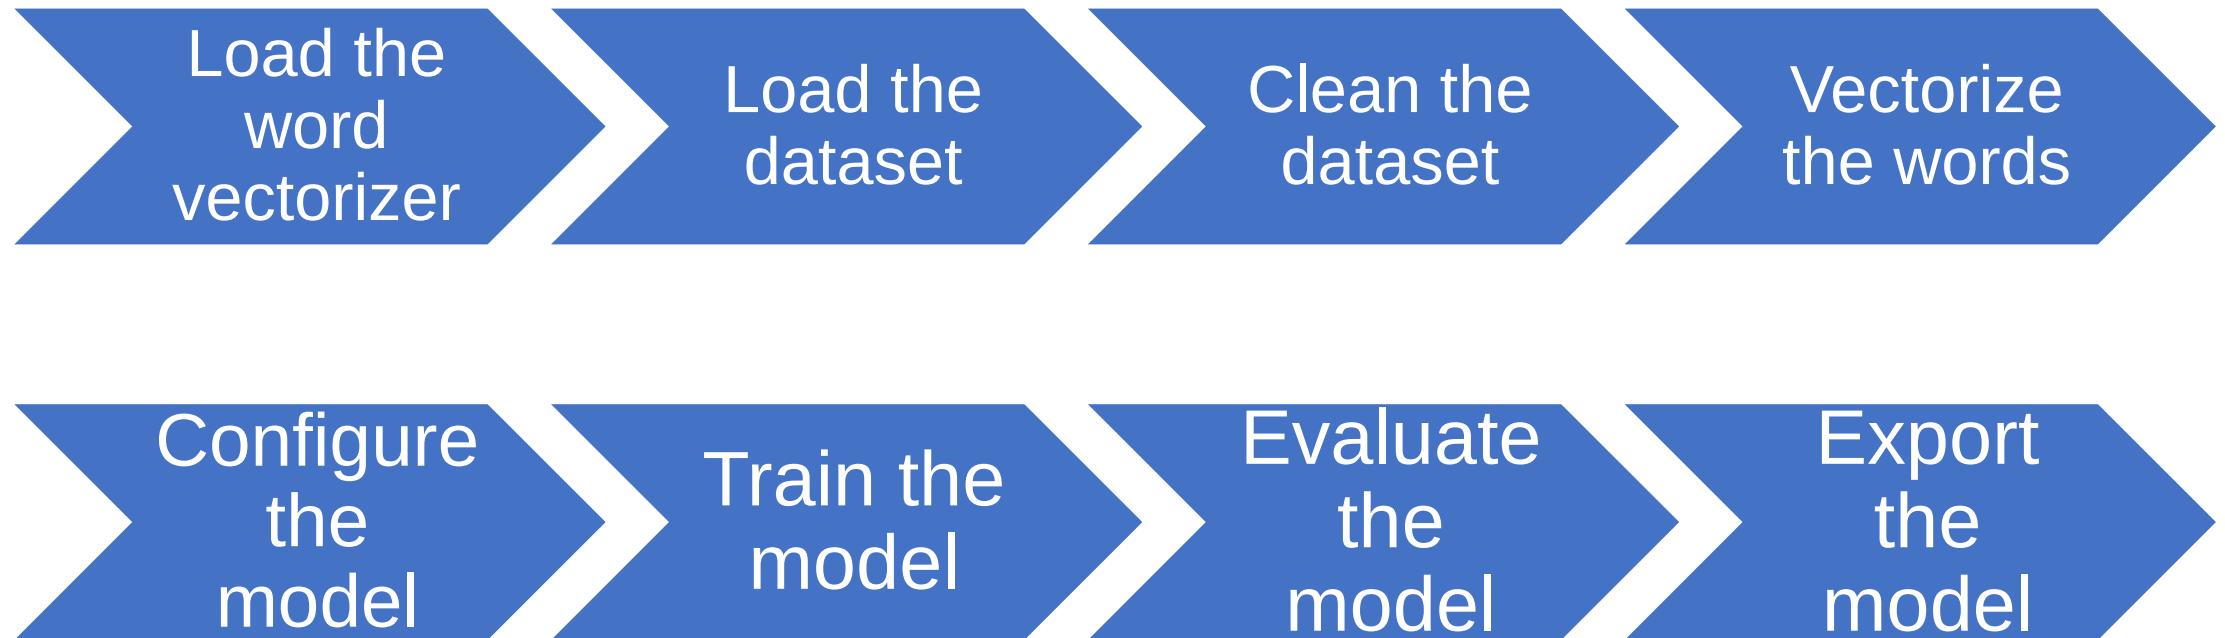
It contains the following 6 fields:

1. target: the polarity of the tweet (0 = negative, 2 = neutral, 4 = positive)

781

#	0	# 1467810369	=	A Mon Apr 06 22:1...	=	A NO_QUERY	=	A _TheSpecialOne_	=	A @switchfoot htt...	=
target		id		date		flag		user		text	
				774362 unique values		1 unique value		659775 unique values		1581465 unique values	
0	0	1467810672		Mon Apr 06 22:19:49 PDT 2009		NO_QUERY		scotthamilton		is upset that he can't update his Facebook by texting it... and might cry as a result School today ...	
0	1467810917			Mon Apr 06 22:19:53 PDT 2009		NO_QUERY		mattycus		@Kenichan I dived many times for the ball. Managed to save 50% The rest go out of bounds	
0	1467811184			Mon Apr 06 22:19:57 PDT 2009		NO_QUERY		ElleCTF		my whole body feels itchy and like its on fire	
0	1467811193			Mon Apr 06 22:19:57 PDT 2009		NO_QUERY		Karoli		@nationwideclass no, it's not behaving at all. i'm mad. why am i here? because I can't see you all o...	
0	1467811372			Mon Apr 06 22:20:00 PDT 2009		NO_QUERY		joy_wolf		@Kwesidei not the i...	

Overview



Vectorizing words

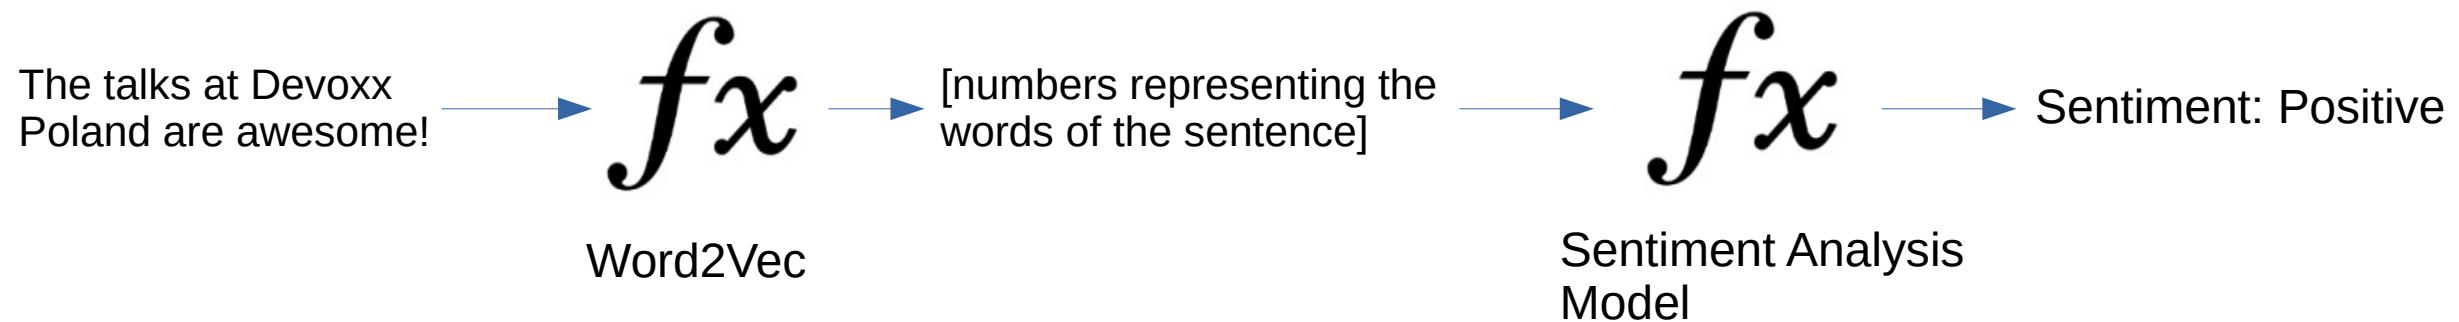
Goal:

Transform the words into vectors (numbers) that the model can understand

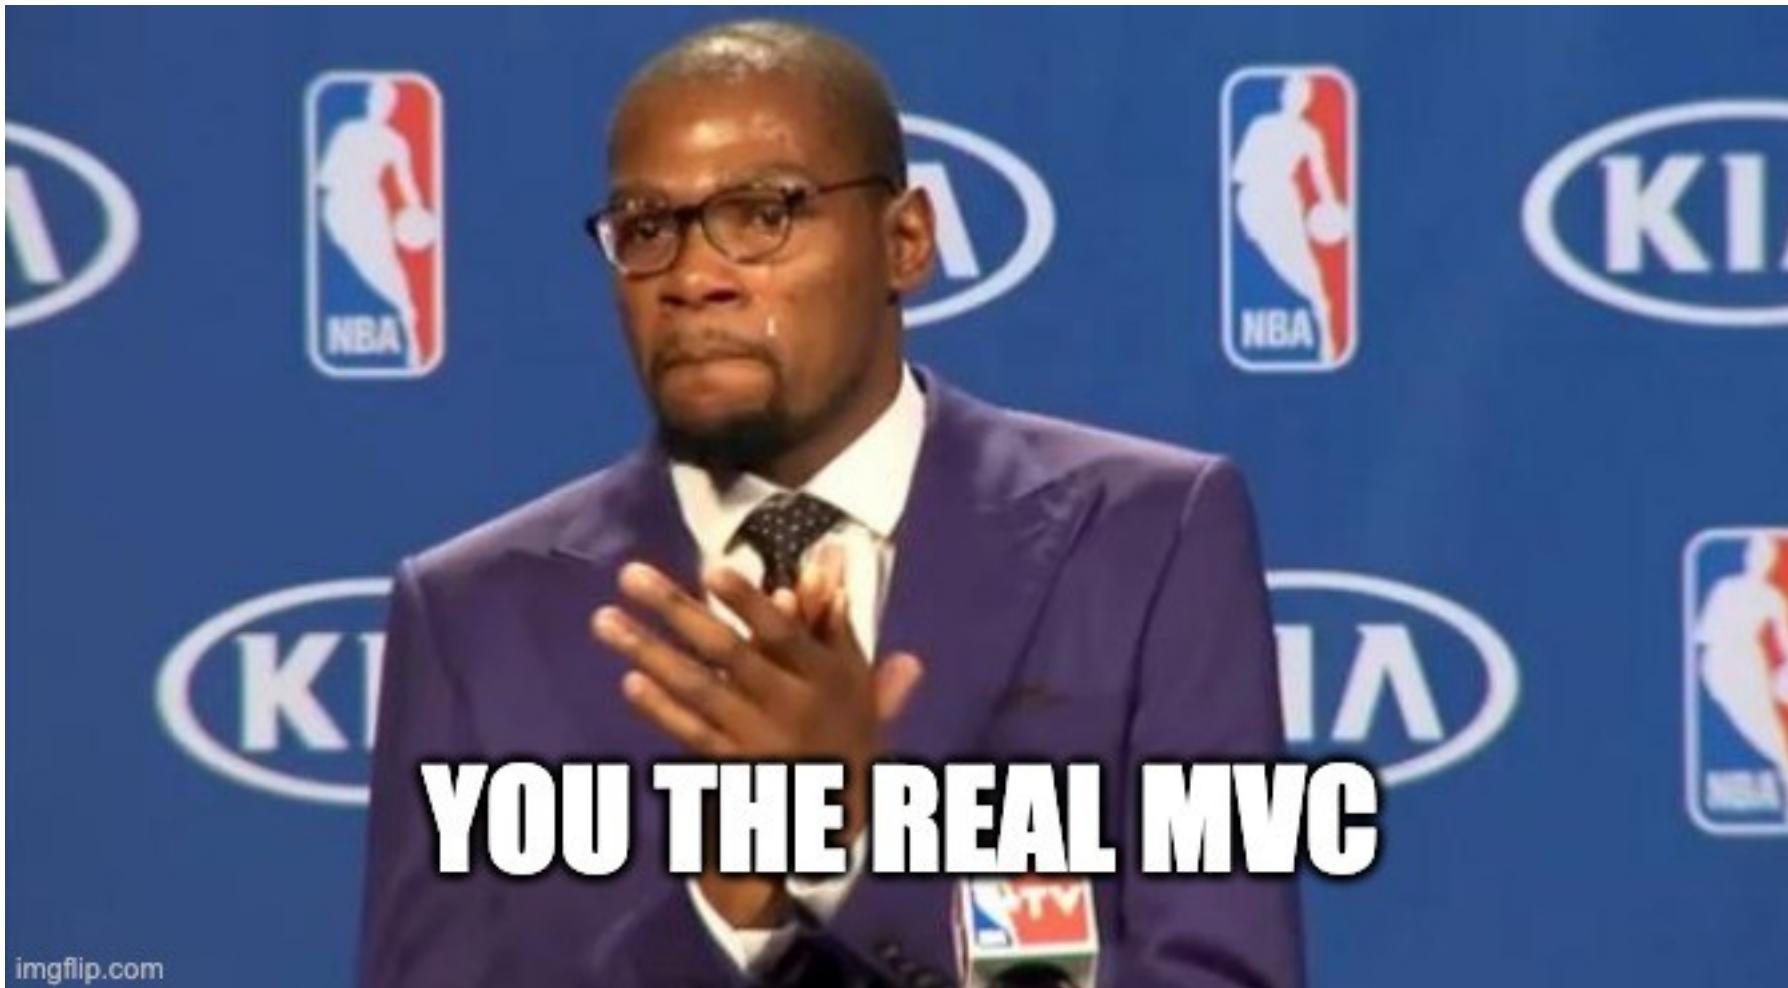
Word2Vec

- Two-layer neural network
- Input: a corpus of text
- Output: a vector space
 - Each unique word in the corpus is assigned a corresponding vector in the space

Vectorizing words



Sentiment Analysis – Most Valuable Code



Load the word vectorizer

Load the dataset

Clean the dataset

Vectorize the words

Configure the model

Train the model

Evaluate the model

Export the model

Goal:

Transform the words into vectors (numbers) that the model can understand



```
WordVectors wordVectors = WordVectorSerializer.loadStaticModel(new File(wordVectorsPath));
```

- This is loading a trained Word2Vec model from Google, can be downloaded from <https://dl4jdata.blob.core.windows.net/resources/wordvectors/GoogleNews-vectors-negative300.bin.gz>
- Example to train your own word2vec model at SimpleExampleWord2Vec.java at <https://github.com/ellerenad/deeplearning4j-playground>

Load the word vectorizer

Load the dataset

Clean the dataset

Vectorize the words

Configure the model

Train the model

Evaluate the model

Export the model



```
LineIterator lineIterator = FileUtils.lineIterator(new File(path), "UTF-8");
int counter = 0;
String line = null;
while (lineIterator.hasNext()) {
    line = lineIterator.nextLine();
    try {
        String[] parts = line.split("\\",\"");
        String label = getLabel(parts);
        String data = getData(parts);
        // Dataset splitting (test and training)
        if (counter % 2 == 0) {
            trainingLabelList.add(label);
            trainingDataList.add(data);
        } else {
            testLabelList.add(label);
            testDataList.add(data);
        }
    } catch(Exception ex){
        log.error("Exception: {}. Counter = {}. Line = {}", ex, counter, line);
    }
    counter++;
}
```

Load the word vectorizer

Load the dataset

Clean the dataset

Vectorize the words

Configure the model

Train the model

Evaluate the model

Export the model



```
LineIterator lineIterator = FileUtils.lineIterator(new File(path), "UTF-8");
int counter = 0;
String line = null;
while (lineIterator.hasNext()) {
    line = lineIterator.nextLine();
    try {
        String[] parts = line.split("\\",\"");
        String label = getLabel(parts);
        String data = getData(parts);
        // Dataset splitting (test and training)
        if (counter % 2 == 0) {
            trainingLabelList.add(label);
            trainingDataList.add(data);
        } else {
            testLabelList.add(label);
            testDataList.add(data);
        }
    } catch(Exception ex){
        log.error("Exception: {}. Counter = {}. Line = {}", ex, counter, line);
    }
    counter++;
}
```

Load the word vectorizer

Load the dataset

Clean the dataset

Vectorize the words

Configure the model

Train the model

Evaluate the model

Export the model



```
LineIterator lineIterator = FileUtils.lineIterator(new File(path), "UTF-8");
int counter = 0;
String line = null;
while (lineIterator.hasNext()) {
    line = lineIterator.nextLine();
    try {
        String[] parts = line.split("\",\"");
        String label = getLabel(parts);
        String data = getData(parts);
        // Dataset splitting (test and training)
        if (counter % 2 == 0) {
            trainingLabelList.add(label);
            trainingDataList.add(data);
        } else {
            testLabelList.add(label);
            testDataList.add(data);
        }
    } catch(Exception ex){
        log.error("Exception: {}. Counter = {}. Line = {}", ex, counter, line);
    }
    counter++;
}
```

Load the word vectorizer

Load the dataset

Clean the dataset

Vectorize the words

Configure the model

Train the model

Evaluate the model

Export the model



```
LineIterator lineIterator = FileUtils.lineIterator(new File(path), "UTF-8");
int counter = 0;
String line = null;
while (lineIterator.hasNext()) {
    line = lineIterator.nextLine();
    try {
        String[] parts = line.split("\\",\"");
        String label = getLabel(parts);
        String data = getData(parts);
        // Dataset splitting (test and training)
        if (counter % 2 == 0) {
            trainingLabelList.add(label);
            trainingDataList.add(data);
        } else {
            testLabelList.add(label);
            testDataList.add(data);
        }
    } catch(Exception ex){
        log.error("Exception: {}. Counter = {}. Line = {}", ex, counter, line);
    }
    counter++;
}
```

Load the word vectorizer

Load the dataset

Clean the dataset

Vectorize the words

Configure the model

Train the model

Evaluate the model

Export the model



```
LineIterator lineIterator = FileUtils.lineIterator(new File(path), "UTF-8");
int counter = 0;
String line = null;
while (lineIterator.hasNext()) {
    line = lineIterator.nextLine();
    try {
        String[] parts = line.split("\\",\"");
        String label = getLabel(parts);
        String data = getData(parts);
        // Dataset splitting (test and training)
        if counter % 2 == 0) {
            trainingLabelList.add(label);
            trainingDataList.add(data);
        } else {
            testLabelList.add(label);
            testDataList.add(data);
        }
    } catch(Exception ex){
        log.error("Exception: {}. Counter = {}. Line = {}", ex, counter, line);
    }
    counter++;
}
```

Load the word vectorizer

Load the dataset

Clean the dataset

Vectorize the words

Configure the model

Train the model

Evaluate the model

Export the model



```
DataWrapper dataWrapper = getCleanData(path);

// The sentence providers shuffle the data internally
LabeledSentenceProvider trainingSentenceProvider = new CollectionLabeledSentenceProvider(
    dataWrapper.getTraining_data(),
    dataWrapper.getTraining_labels(),
    rng);

// The iterator vectorizes internally the words
DataSetIterator trainingDataSetIterator = new CnnSentenceDataSetIterator.Builder(Format.CNN2D)
    .sentenceProvider(trainingSentenceProvider)
    .wordVectors(wordVectors) ←
    .minibatchSize(minibatchSize)
    .maxSentenceLength(maxSentenceLength)
    .useNormalizedWordVectors(false)
    .build();
```

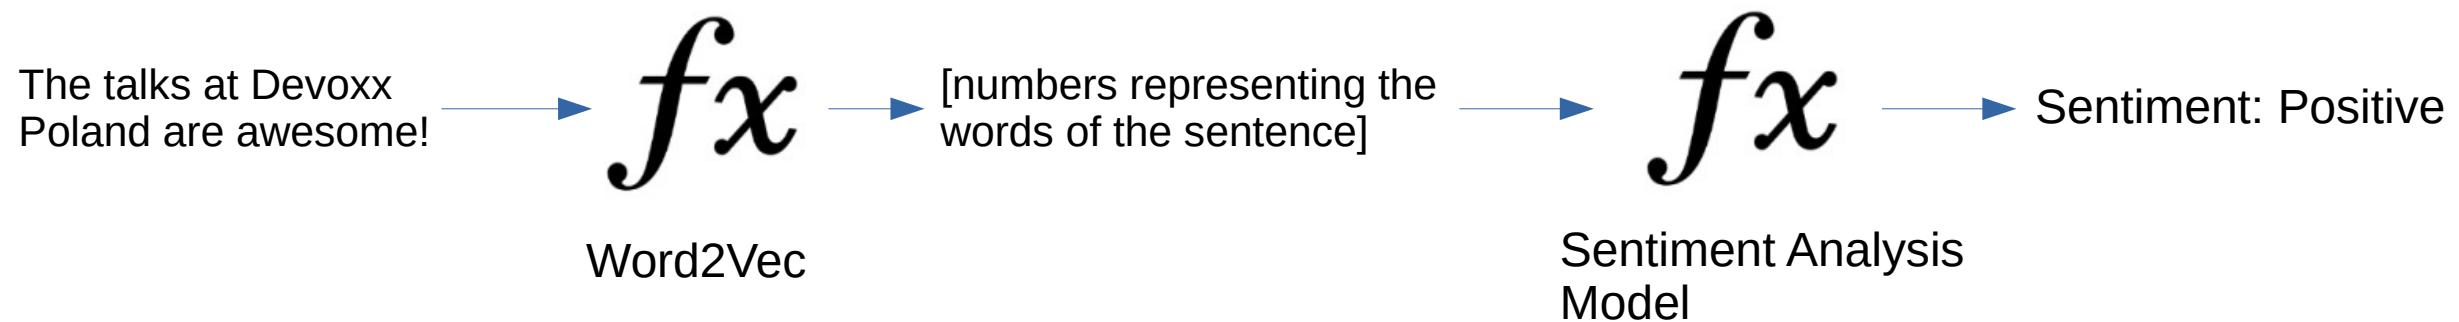
The diagram shows a sequence of eight blue chevron-shaped arrows pointing from left to right, representing a workflow. The fourth arrow from the left is orange and labeled 'Vectorize the words'. Two orange arrows point downwards from this orange chevron to the corresponding code lines in the Java snippet: one points to 'wordVectors(wordVectors)' and another points to '.build()'.



Configure the model

```
ComputationGraphConfiguration config = new NeuralNetConfiguration.Builder()
    .weightInit(WeightInit.RELU)
    .activation(Activation.LEAKYRELU)
    .updater(new Adam(0.01))
    .convolutionMode(ConvolutionMode.Same)
    .l2(0.0001)
    .graphBuilder()
    .addInputs("input")
    .addLayer("cnn3", new ConvolutionLayer.Builder()
        .kernelSize(3, vectorSize)
        .stride(1, vectorSize)
        .nOut(cnnLayerFeatureMaps)
        .build(), "input")
    .addLayer("cnn4", new ConvolutionLayer.Builder()
        .kernelSize(4, vectorSize)
        .stride(1, vectorSize)
        .nOut(cnnLayerFeatureMaps)
        .build(), "input")
    .addLayer("cnn5", new ConvolutionLayer.Builder()
        .kernelSize(5, vectorSize)
        .stride(1, vectorSize)
        .nOut(cnnLayerFeatureMaps)
        .build(), "input")
    .addVertex("merge", new MergeVertex(), "cnn3", "cnn4", "cnn5")
    .addLayer("globalPool", new GlobalPoolingLayer.Builder()
        .poolingType(globalPoolingType)
        .dropOut(0.5)
        .build(), "merge")
    .addLayer("out", new OutputLayer.Builder()
        .lossFunction(LossFunctions.LossFunction.MCXENT)
        .activation(Activation.SOFTMAX)
        .nOut(2) //2 classes: positive or negative
        .build(), "globalPool")
    .setOutputs("out")
    //Input has shape [minibatch, channels=1, length=1 to 256, 300]
    .setInputTypes(InputType.convolutional(truncateReviewsToLength, vectorSize, 1))
    .bullet();
```

Vectorizing words





Configure the model

```
ComputationGraphConfiguration config = new NeuralNetConfiguration.Builder()
    .weightInit(WeightInit.RELU)
    .activation(Activation.LEAKYRELU)
    .updater(new Adam(0.01))
    .convolutionMode(ConvolutionMode.Same)
    .l2(0.0001)
    .graphBuilder()
    .addInputs("input")
    .addLayer("cnn3", new ConvolutionLayer.Builder()
        .kernelSize(3, vectorSize)
        .stride(1, vectorSize)
        .nOut(cnnLayerFeatureMaps)
        .build(), "input")
    .addLayer("cnn4", new ConvolutionLayer.Builder()
        .kernelSize(4, vectorSize)
        .stride(1, vectorSize)
        .nOut(cnnLayerFeatureMaps)
        .build(), "input")
    .addLayer("cnn5", new ConvolutionLayer.Builder()
        .kernelSize(5, vectorSize)
        .stride(1, vectorSize)
        .nOut(cnnLayerFeatureMaps)
        .build(), "input")
    .addVertex("merge", new MergeVertex(), "cnn3", "cnn4", "cnn5")
    .addLayer("globalPool", new GlobalPoolingLayer.Builder()
        .poolingType(globalPoolingType)
        .dropOut(0.5)
        .build(), "merge")
    .addLayer("out", new OutputLayer.Builder()
        .lossFunction(LossFunctions.LossFunction.MCXENT)
        .activation(Activation.SOFTMAX)
        .nOut(2) //2 classes: positive or negative
        .build(), "globalPool")
    .setOutputs("out")
    //Input has shape [minibatch, channels=1, length=1 to 256, 300]
    .setInputTypes(InputType.convolutional(truncateReviewsToLength, vectorSize, 1))
    .build();
```

Load the word vectorizer

Load the dataset

Clean the dataset

Vectorize the words

Configure the model

Train the model

Evaluate the model

Export the model



```
DataSetIterator trainIter = dataSetIteratorWrapper.getTrain();
DataSetIterator testIter = dataSetIteratorWrapper.getTest();

// Listeners to print information during the training and the final evaluation
net.setListeners(new ScoreIterationListener(100),
                  new EvaluativeListener(testIter, 1, InvocationType.EPOCH_END));
// Train the neural network
net.fit(trainIter, nEpochs);

// Export the neural network
net.save(new File("/example/path/"));
```

Sentiment analysis prediction example

“this code is driving me crazy! what was this guy thinking!? *Annotates the code* oh... it was me”

$P(\text{Negative}) = 0.72$

$P(\text{Positive}) = 0.28$

Label: Negative

Problems during development

- Slow feedback
- Memory management
- Loading the dataset

Sampling for faster feedback



```
head -1000 training.1600000.processed.noemoticon.csv >> training_reduced.csv  
tail -1000 training.1600000.processed.noemoticon.csv >> training_reduced.csv
```

Comparing results

Samples: 800000		Samples: 2000	
Accuracy	0.7472	Accuracy	0.6721
Precision	0.7200	Precision	0.6382
Recall	0.8065	Recall	0.7886
F1 Score	0.7608	F1 Score	0.7055
Execution Time	1581209 ms ~26.3 min	Execution Time	165002 ms ~2.75 min

The diagram illustrates the relationship between the number of samples and execution time. An orange arrow points from the 'Samples: 2000' row to the 'Samples: 800000' row, labeled '40x'. Another orange arrow points from the 'Execution Time' row for 800000 samples (~26.3 min) to the 'Execution Time' row for 2000 samples (~2.75 min), labeled '10x'.

Comparing results

Samples: 2000

Accuracy	0.6721
Precision	0.6382
Recall	0.7886
F1 Score	0.7055
Execution Time	165002 ms ~2.75 min

Loading the word vectors:

154804 ms

~2.58 min



```
WordVectors wordVectors = WordVectorSerializer.loadStaticModel(new File(wordVectorsPath));
```

Memory Management

- Tweak the JVM with the params:



```
-Xms1024m // Initial memory allocation pool  
-Xmx10g // Maximum memory allocation pool
```

More details at: <https://deeplearning4j.konduit.ai/config/config-memory>

Load the data set

On files bigger than 40,000 lines, the `java.util.Scanner` would read just up to the half of the file, messing with the training.



```
// Load dataset  
LineIterator lineIterator = FileUtils.lineIterator(new File(path), "UTF-8");
```

Thanks!

Questions?

All the code can be found at

<https://github.com/ellerenad/deeplearning4j-playground>

Twitter @ellerenad

Github @ellerenad

Blog posts about this talk at
ienjoyssoftware.dev

Scan the QR code to get all
the relevant links!

